



MULTIPLE LABELS

Deep Learning for Computer Vision

Arthur Douillard

Few-Shot Learning



Few-Shots Learning

LFW: Labeled Face in the Wild



Omniglot



Kaggle's Humpback Whale identification

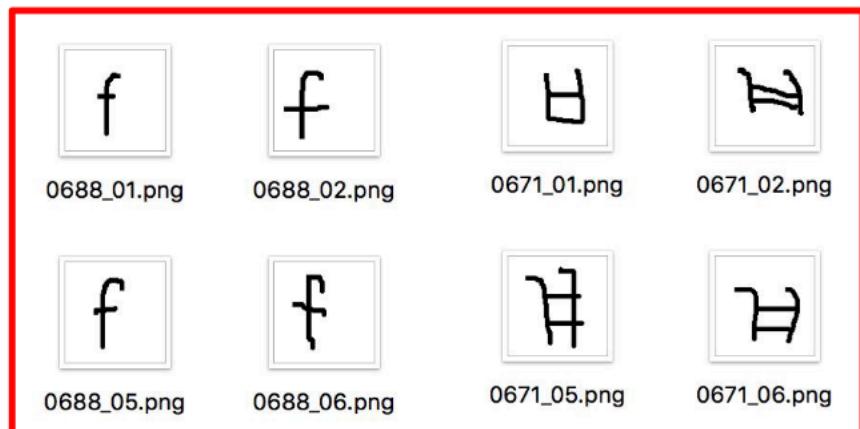


Few-Shots Learning



1. Learn to classify the few labeled samples in the **background set**
2. With a few labeled samples in **support set**, classify the **query set**

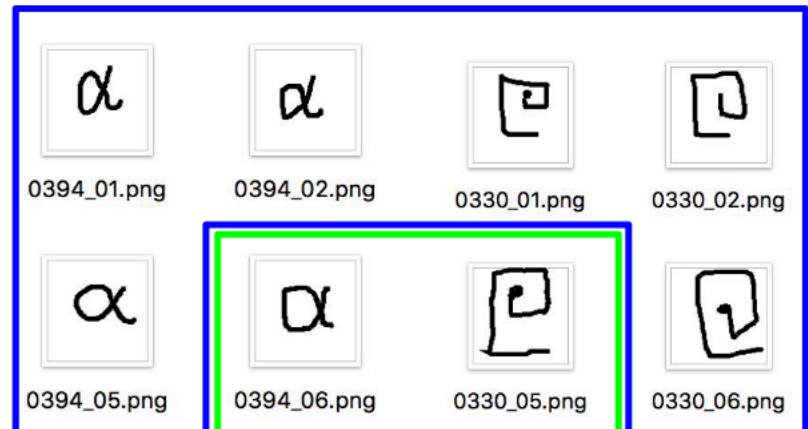
Training:



Latin F

Korean B/P

Testing:



Greek Alpha

Futurama F

Number of labeled samples / class in background set: **K-shots**

Number of classes in query set: **N-ways**



A huge, potentially growing number of classes.

Less than a dozen labeled samples per class.

Discriminative model is impossible.

What if we learn a metric instead?



Distance:

$$d(x_1, x_2) = \|f(x_1) - f(x_2)\|_2 \in \mathbb{R}^+$$

Similarity:

$$s(x_1, x_2) = \cos(f(x_1), f(x_2)) \in [-1, +1]$$

With $f(x) \in \mathbb{R}^d$ a features extractor (e.g. ConvNet).

Given two images of the class, we want:

- Minimize distance
- Maximize similarity

Given two images of different classes, we want:

- Maximize distance
- Minimize similarity

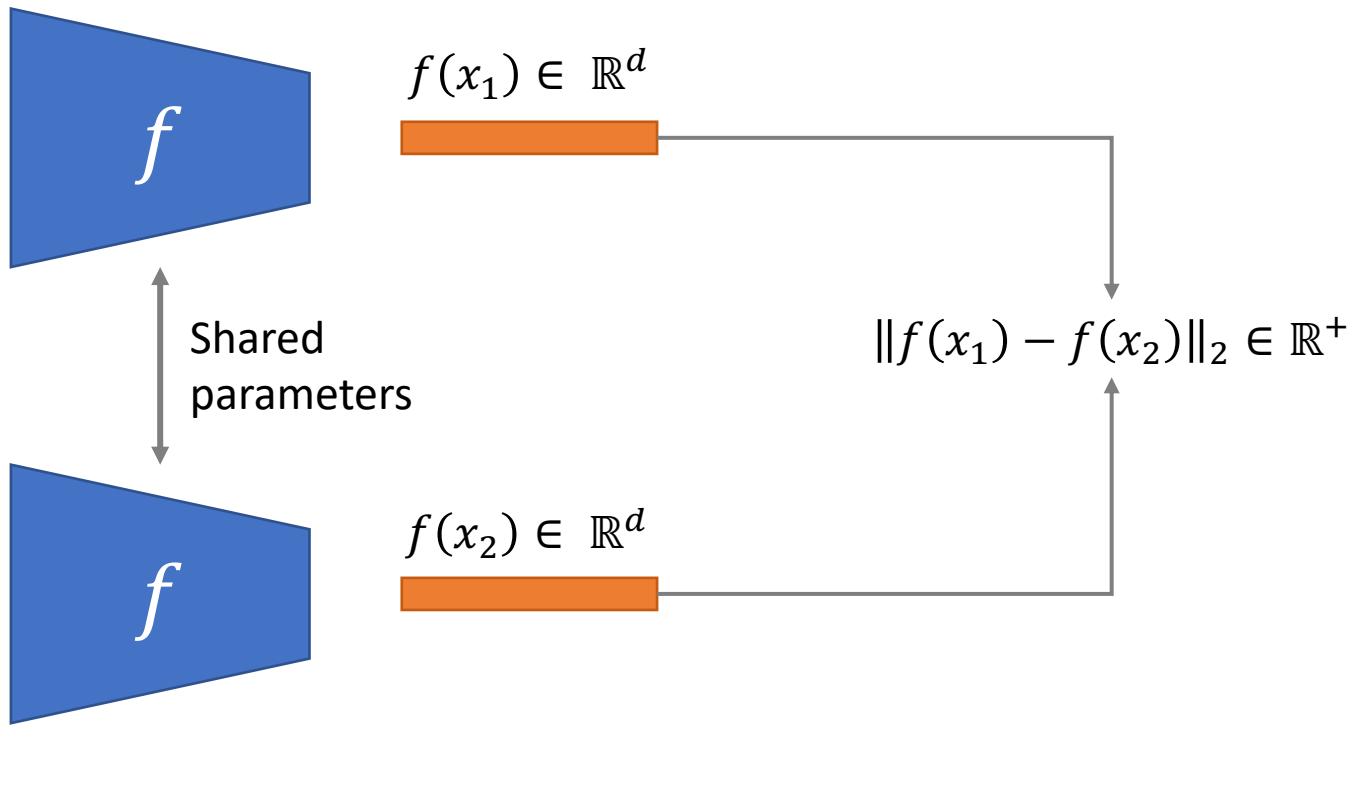
Siamese Network



x_1



x_2

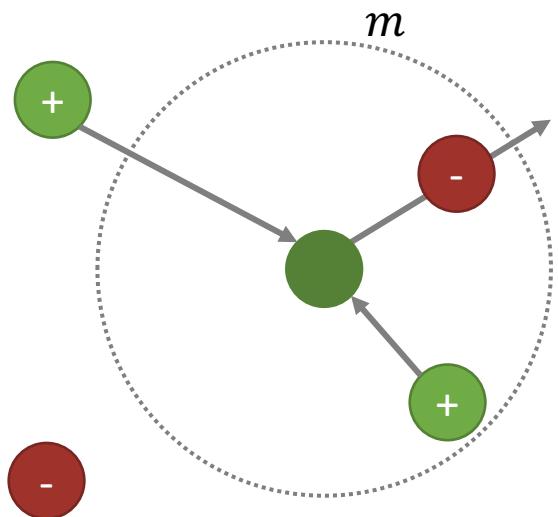


Contrastive Loss



$$D = \|f(x_1) - f(x_2)\|_2 \in \mathbb{R}^+$$

$$\mathcal{L}_{contrastive}(y, D) = \frac{1}{2}(1 - y)D^2 + \frac{1}{2}y \max(m - D, 0)^2$$



Weaknesses of Pairwise Models



- The margin m may be hard to tune, especially because distributions can change through training
- A double-margin may improve to avoid collapsing all positive samples together
- Try to learn an absolute distance between images

Triplet Network



We want to learn **relative distance** between samples

Given an anchor x_a , we want to have a small distance with a positive (same class) x_+ :

$$\min \|f(x_a) - f(x_+)\|_2$$

And maximize with a negative (different class) x_- :

$$\max \|f(x_a) - f(x_-)\|_2$$

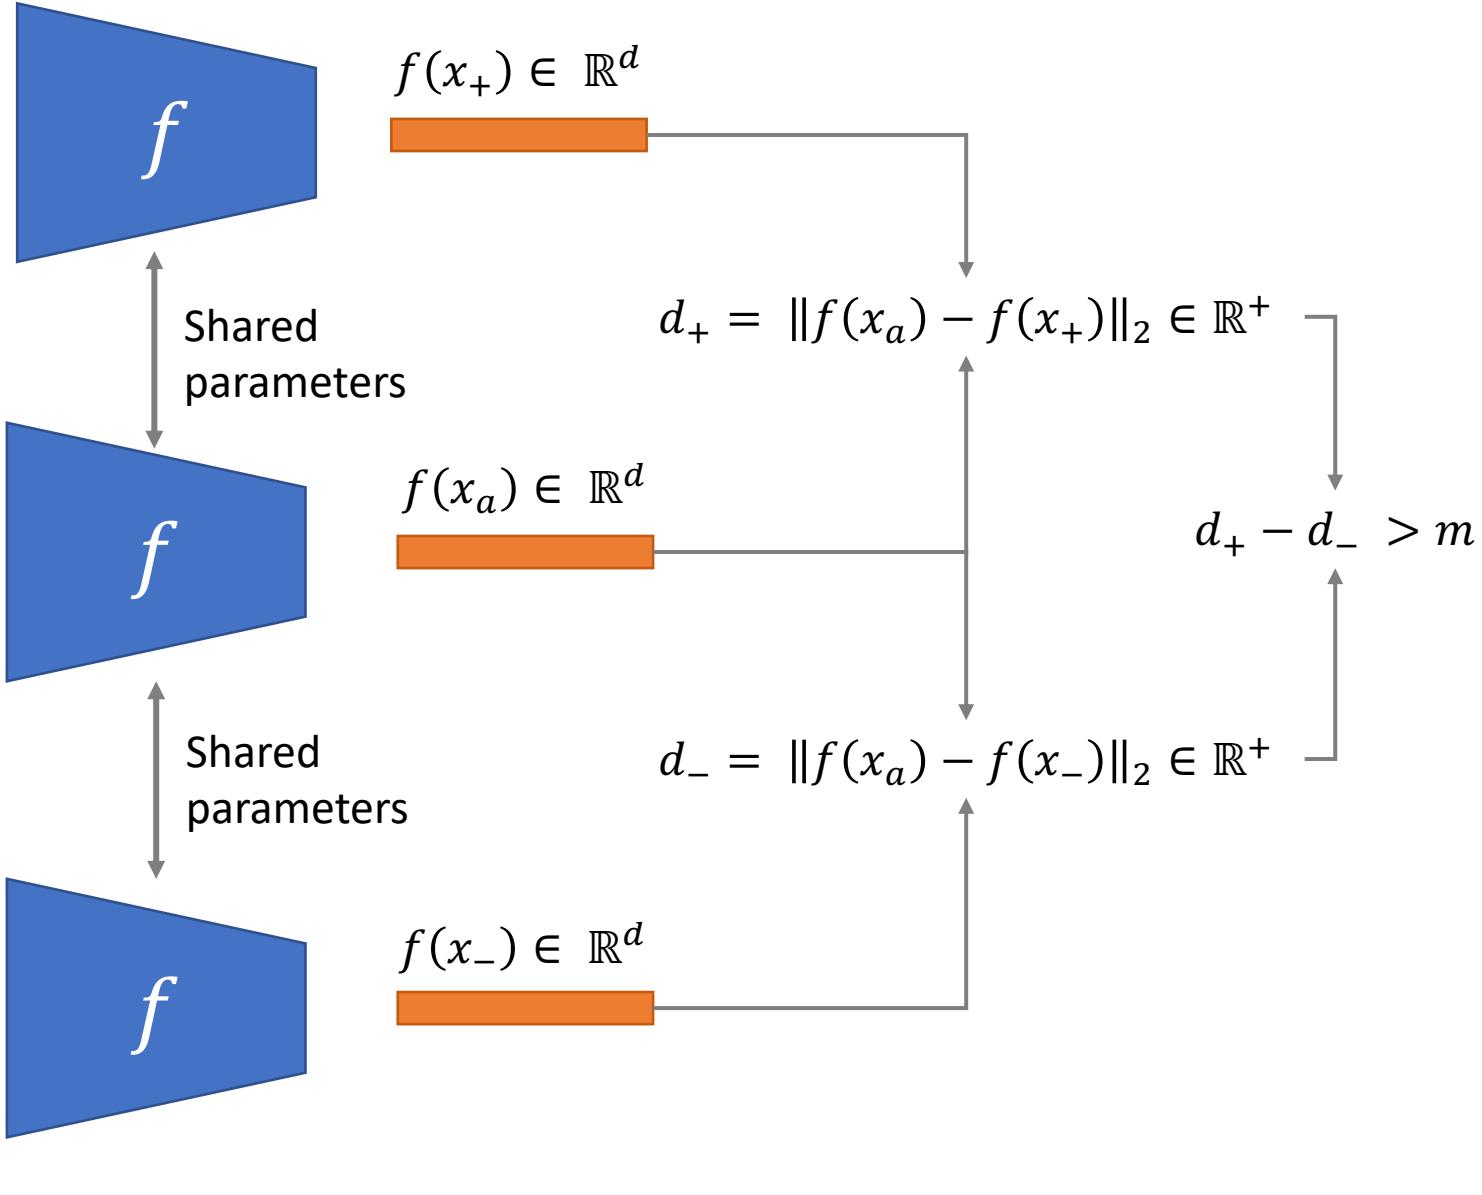
Therefore:

$$\begin{aligned} \|f(x_a) - f(x_+)\|_2 &> \|f(x_a) - f(x_-)\|_2 \\ \|f(x_a) - f(x_+)\|_2 - \|f(x_a) - f(x_-)\|_2 &> 0 \end{aligned}$$

Add a margin m to ensure extra separability:

$$\|f(x_a) - f(x_+)\|_2 - \|f(x_a) - f(x_-)\|_2 > m$$

Triplet Network

 x_+  x_a  x_- 

Hard Negative Mining



Most triplets are easy.

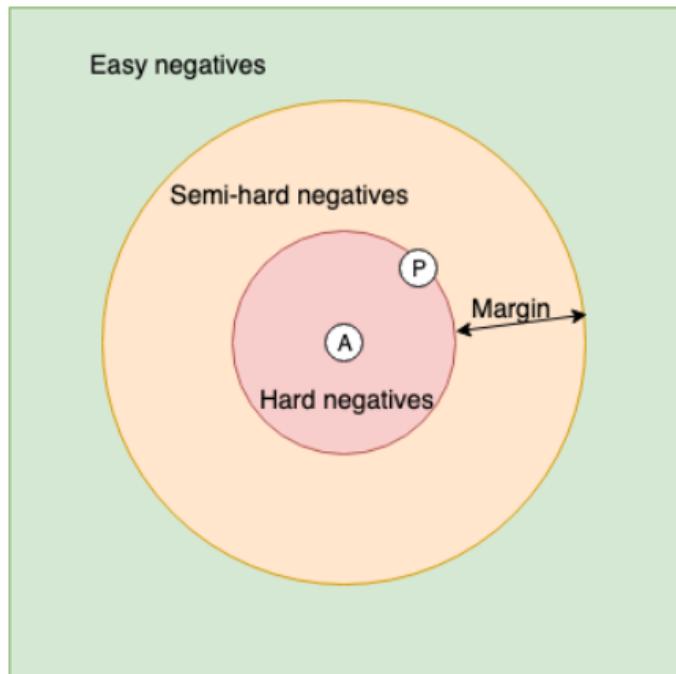
We want to sample either:

Hard negatives:

$$\|f(x_a) - f(x_+)\|_2 > \|f(x_a) - f(x_-)\| + m$$

Semi-Hard negatives:

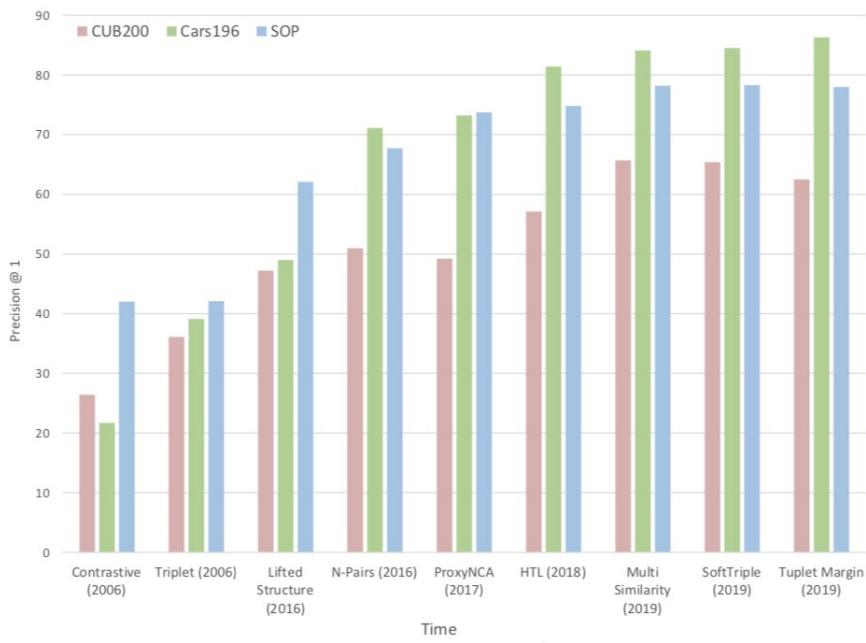
$$\|f(x_a) - f(x_+)\|_2 > \|f(x_a) - f(x_-)\|$$



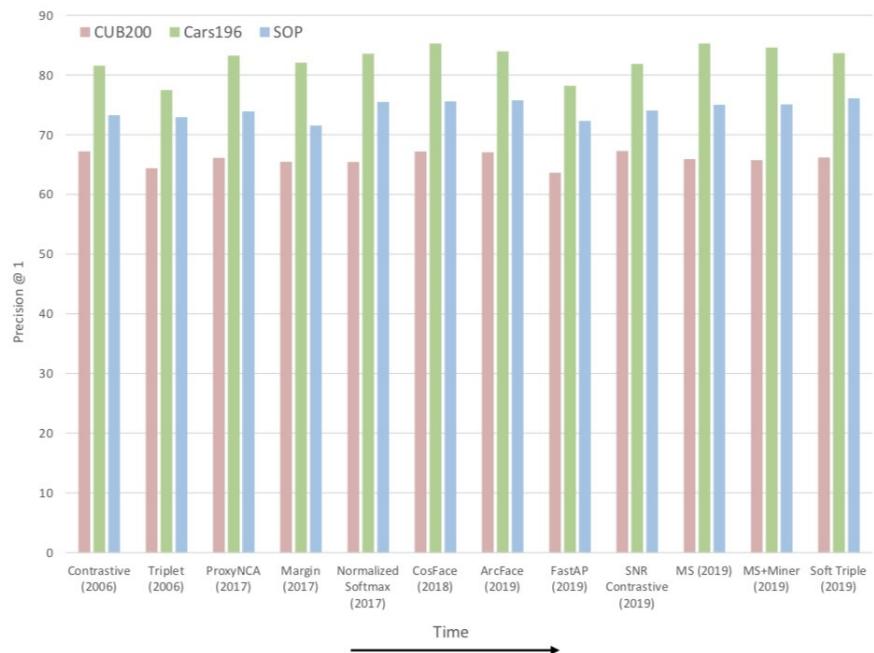
Pretty much essential to have SotA performance!



Troubles in Metric Learning...



(a) The trend according to papers



(b) The trend according to reality

Fig. 4. Papers versus Reality: the trend of Precision@1 of various methods over the years.

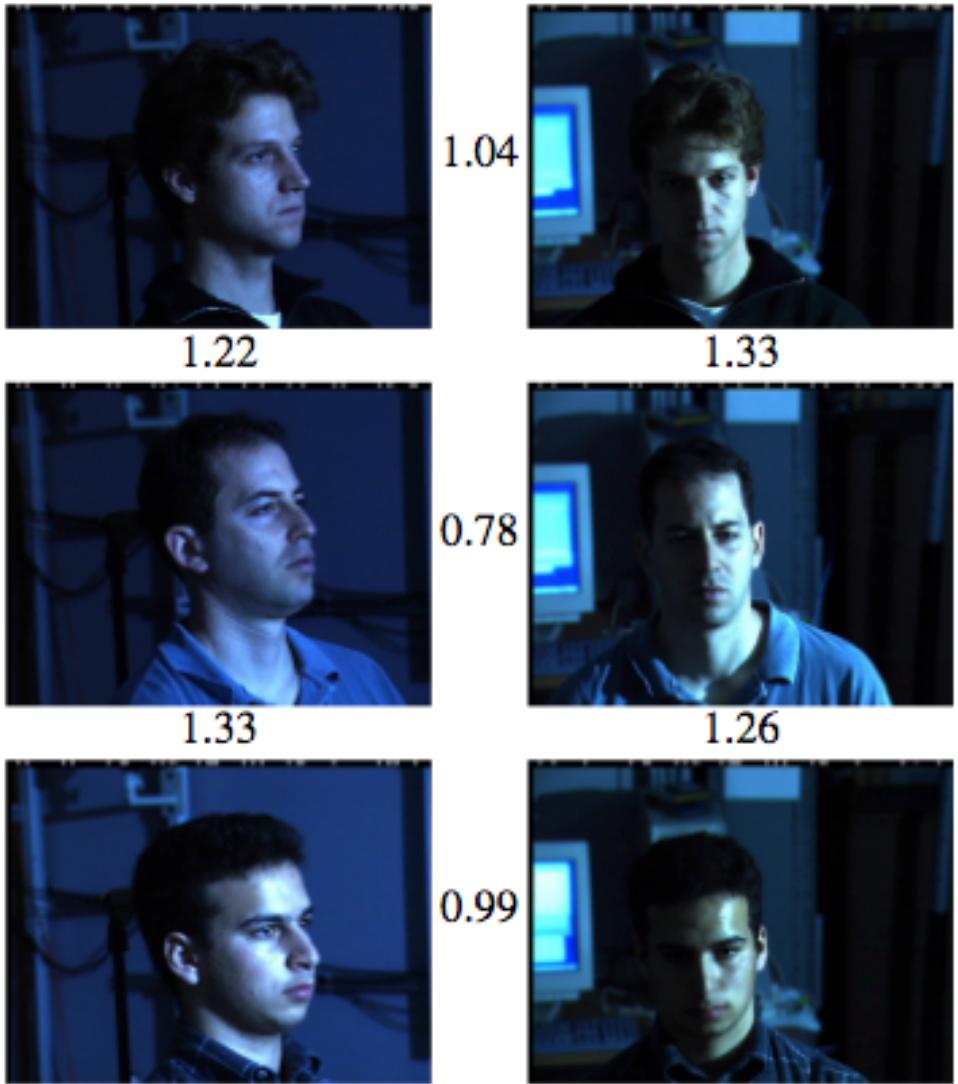
The gain of more recent Metric Learning models may come from:

- Better backbone
- Better hyperparameters tuning
- Better data augmentation



Triplet Network with semi-hard negative mining.

Pretty much solved the LFW dataset.



Distance between pairs

Meta-Learning MAML



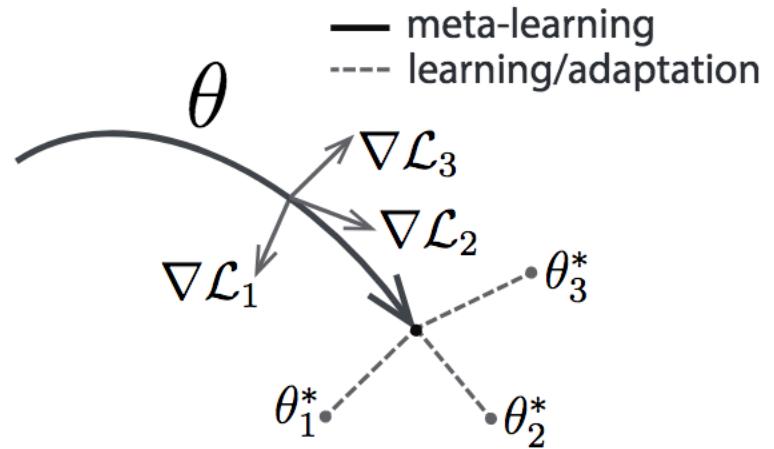
Learning to learn:

→ Given a few samples can learn faster

Outer and inner loop:

- Inner loop learns to classify well a few labeled samples.
- Outer loop learns to have good weights for the inner loop.

During inference, perform only inner loop.



Algorithm 1 Model-Agnostic Meta-Learning

Require: $p(\mathcal{T})$: distribution over tasks

Require: α, β : step size hyperparameters

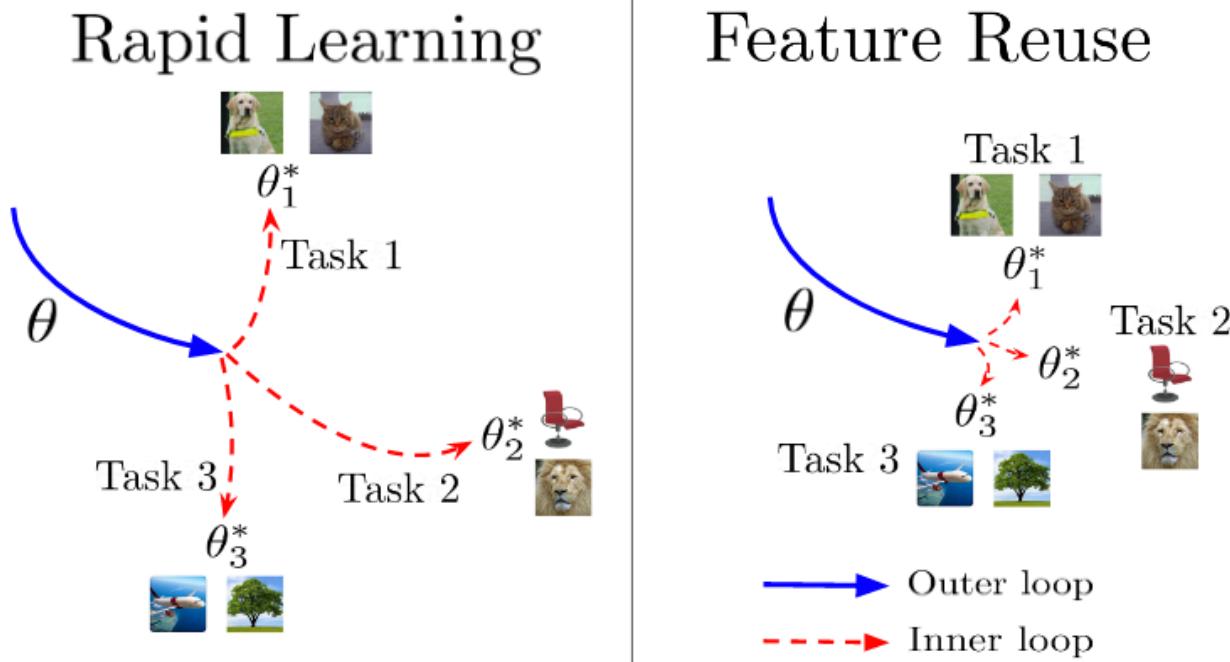
- 1: randomly initialize θ
- 2: **while** not done **do**
- 3: Sample batch of tasks $\mathcal{T}_i \sim p(\mathcal{T})$
- 4: **for all** \mathcal{T}_i **do**
- 5: Evaluate $\nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$ with respect to K examples
- 6: Compute adapted parameters with gradient descent: $\theta'_i = \theta - \alpha \nabla_{\theta} \mathcal{L}_{\mathcal{T}_i}(f_{\theta})$
- 7: **end for**
- 8: Update $\theta \leftarrow \theta - \beta \nabla_{\theta} \sum_{\mathcal{T}_i \sim p(\mathcal{T})} \mathcal{L}_{\mathcal{T}_i}(f_{\theta'_i})$
- 9: **end while**

Meta-Learning MAML



Does MAML *really* learns to learn (**rapid learning**)?

More probably it manages to learn features that generalize well (**feature reuse**)



Self-Supervised Learning

(also Unsupervised Learning)

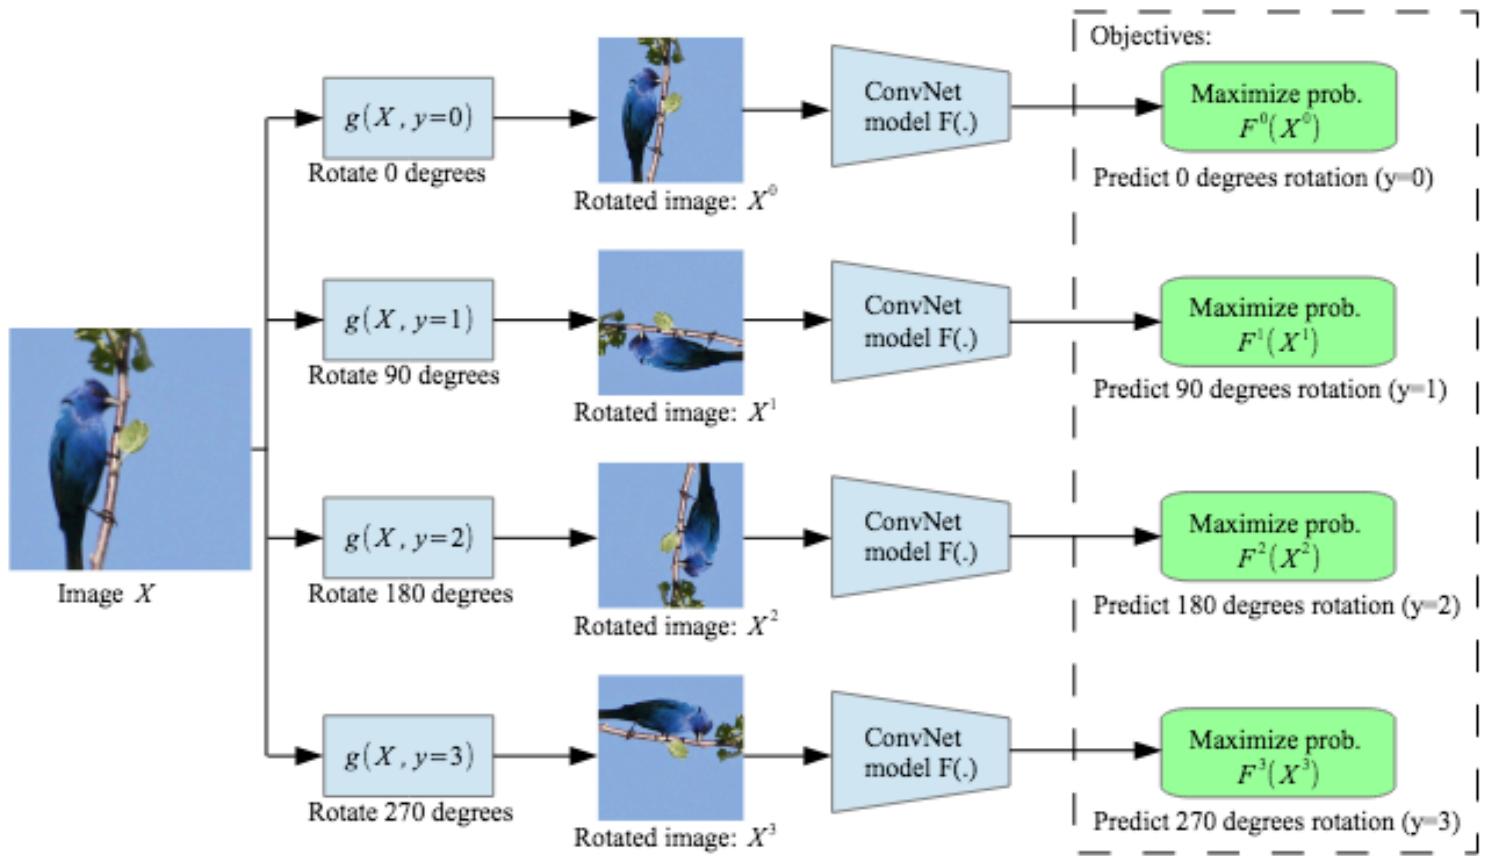


Setting

1. Learn a backbone with a lot of **unlabeled data** using a self-designed objective
 2. Freeze the backbone
 3. Learn a classifier on top of it on labeled data.
- Evaluate how much the learned representation (features extracted by ConvNet) is useful.

It can be useful to have a good model pretraining in order to do transfer learning later.

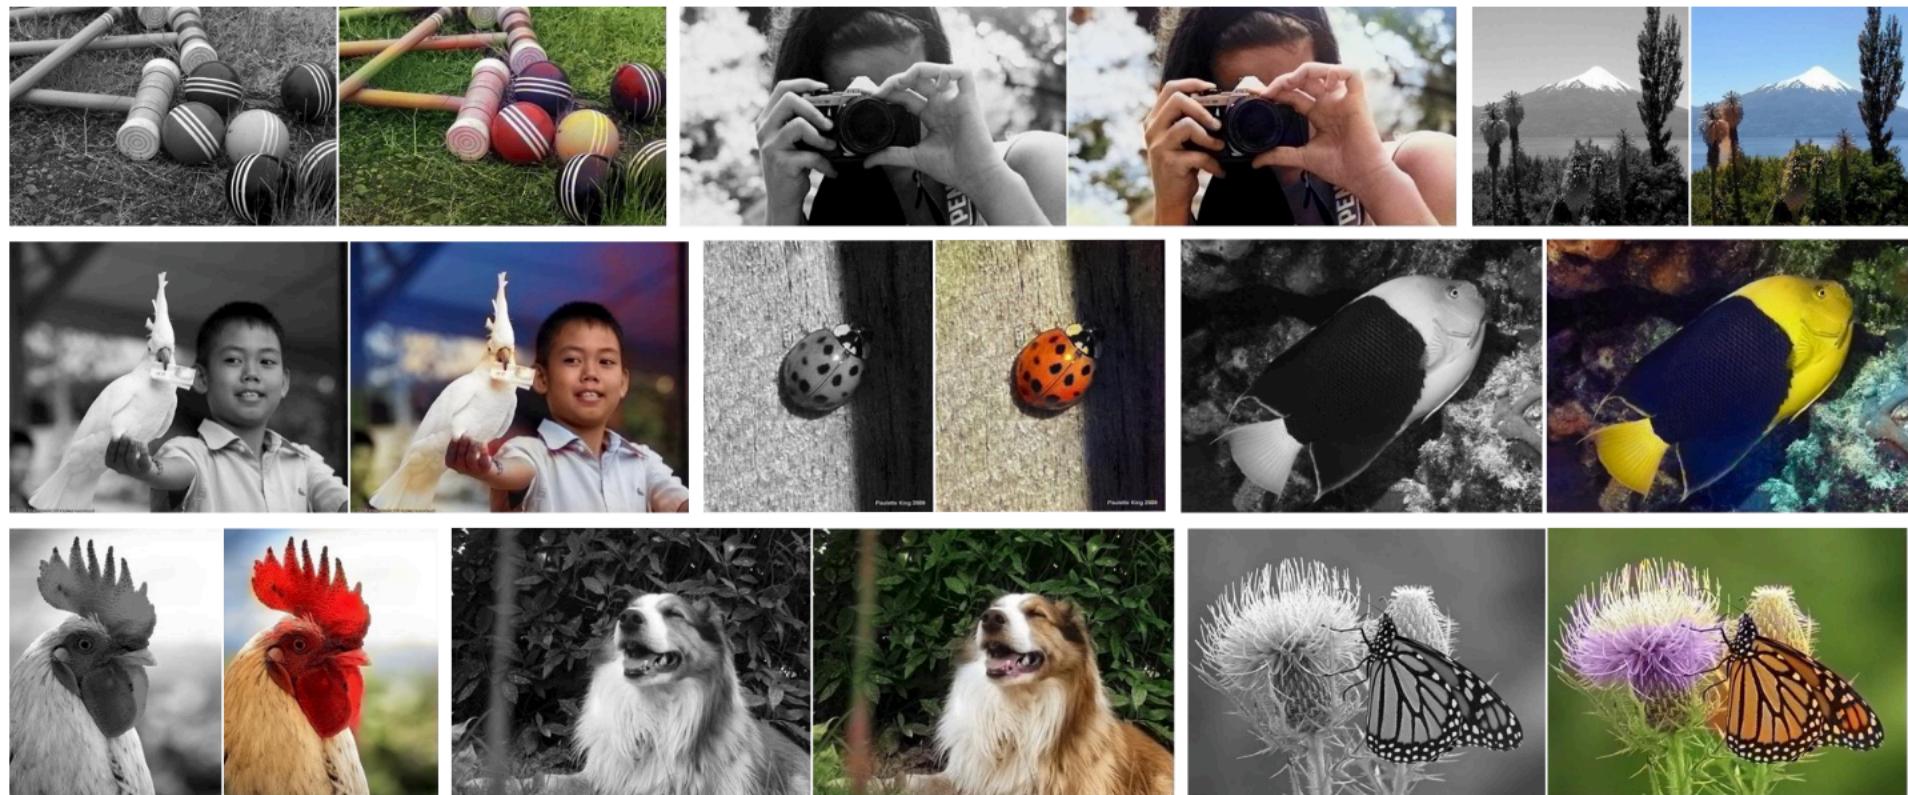
Rotation



Learn what is the rotation applied.

→ Thus it must learn what is the structure of the visual world.

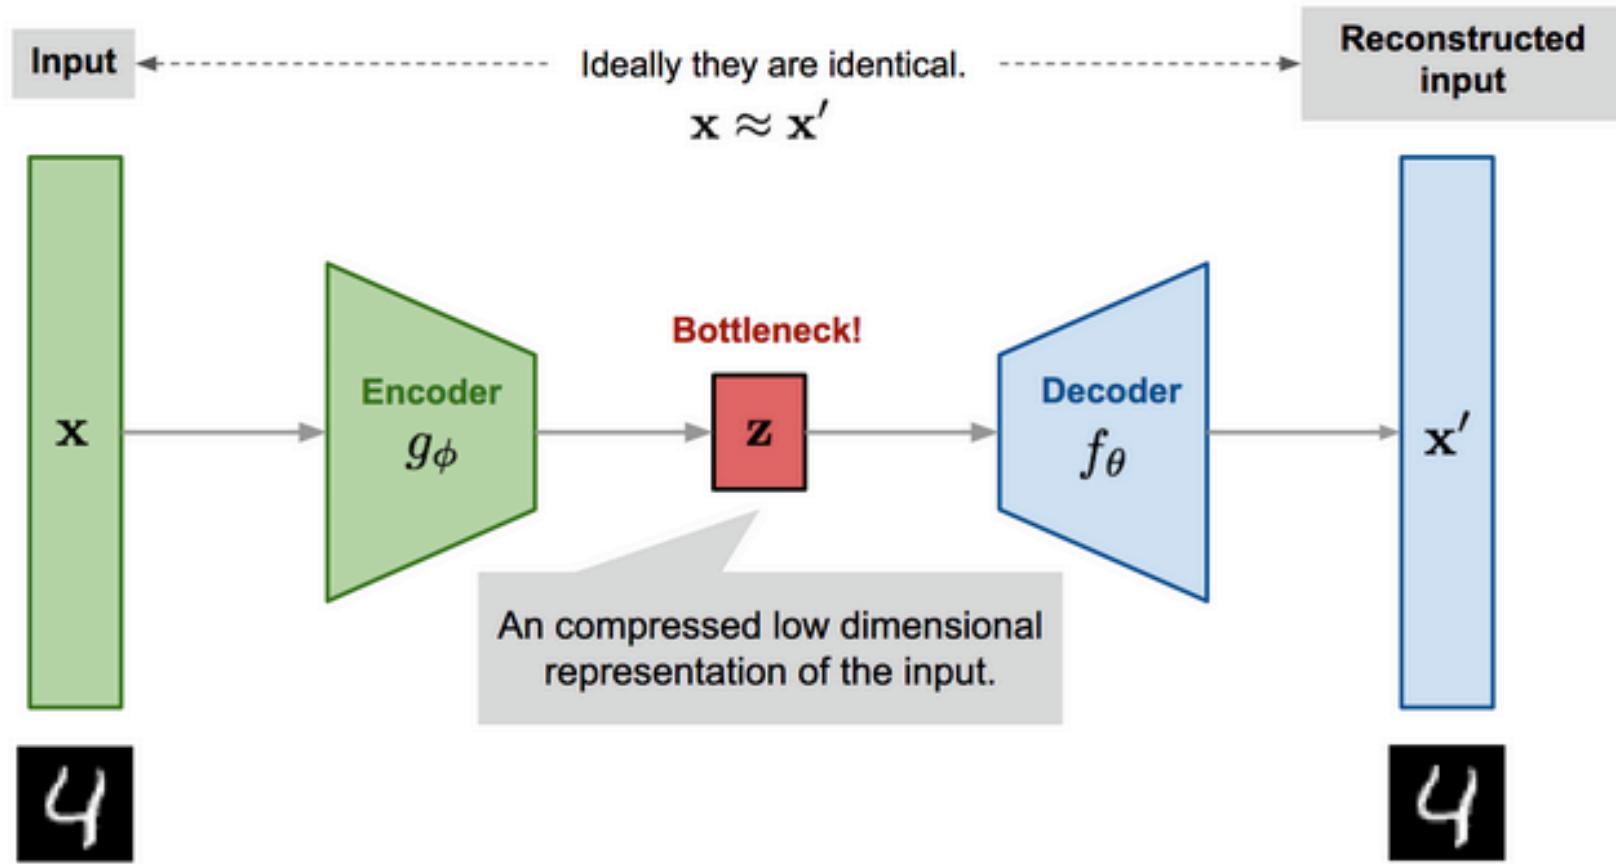
Colorization



Learns to predict the colors from “*grayscale*” images.

- Regression problem
- Done in the LAB space instead of RGB

Auto-Encoder



Compress an image and then reconstruct it.

- Similarly to method of dimensionality reduction like PCA
- It must only keep important features.

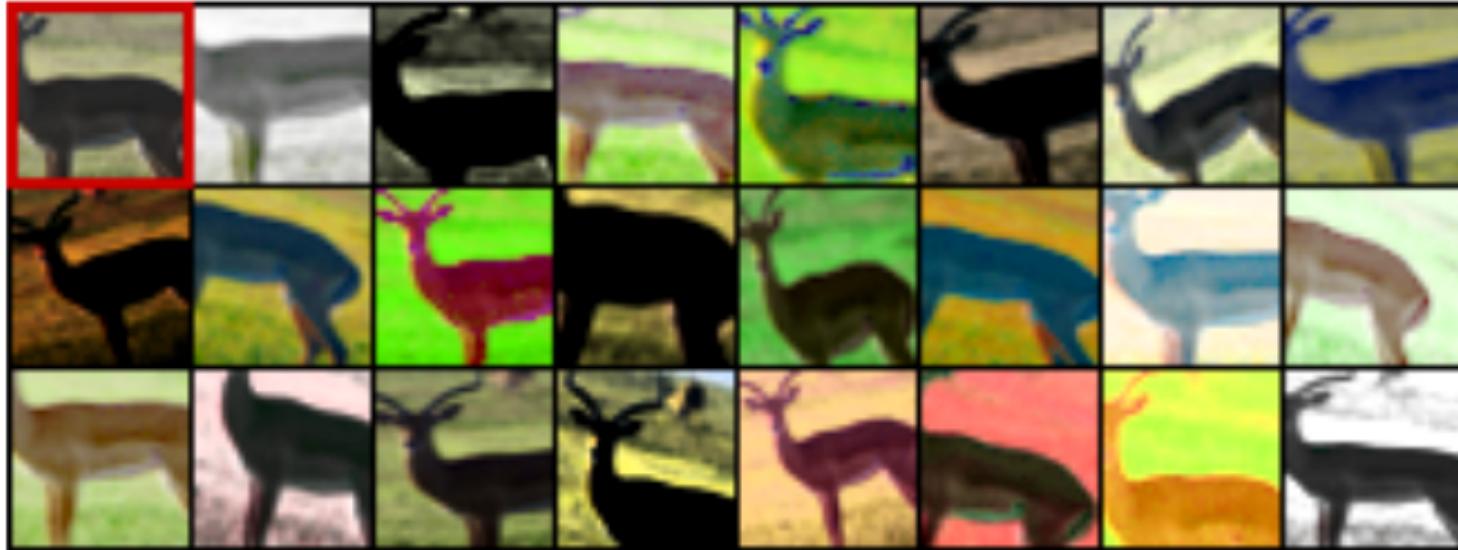
Constrastive



Each image is considered as a class.

New samples of this “class” are generated with heavy data augmentation.

Trained with usual softmax + cross-entropy.

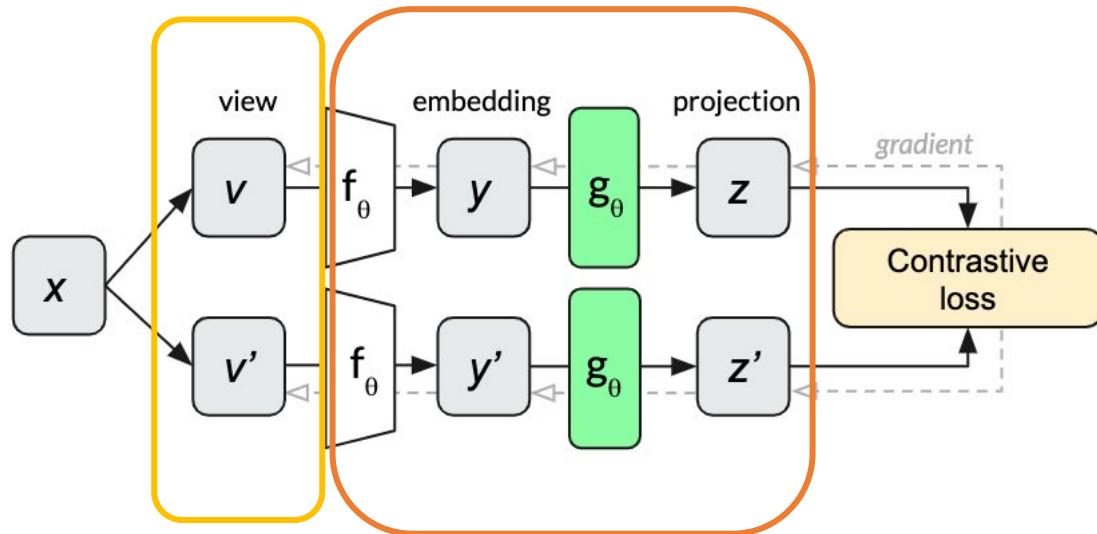


Most of future self-supervised models also try to:

- Bring together the same image augmented differently
- Push away all others images



Batch of images is augmented twice differently



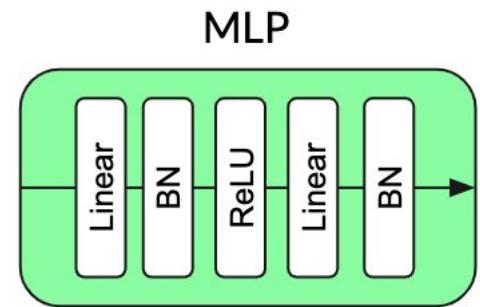
Extract features with ConvNet f_θ ,
then project it with a small MLP
to produce $z \in \mathbb{R}^d$

Alternative version of the contrastive loss.

→ Bring together the same image augmented differently

→ Push away all others images of the batch

$$-\log \frac{\exp(s_{i,j}/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(s_{i,k}/\tau)}$$





Alternative version of the contrastive loss.

- Bring together the same image augmented differently
- Push away all others images of the batch

$$-\log \frac{\exp(s_{i,j}/\tau)}{\sum_{k=1}^{2N} \mathbb{1}_{[k \neq i]} \exp(s_{i,k}/\tau)}$$

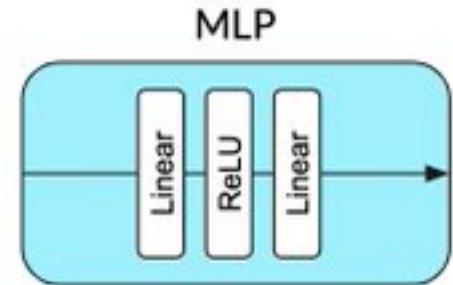
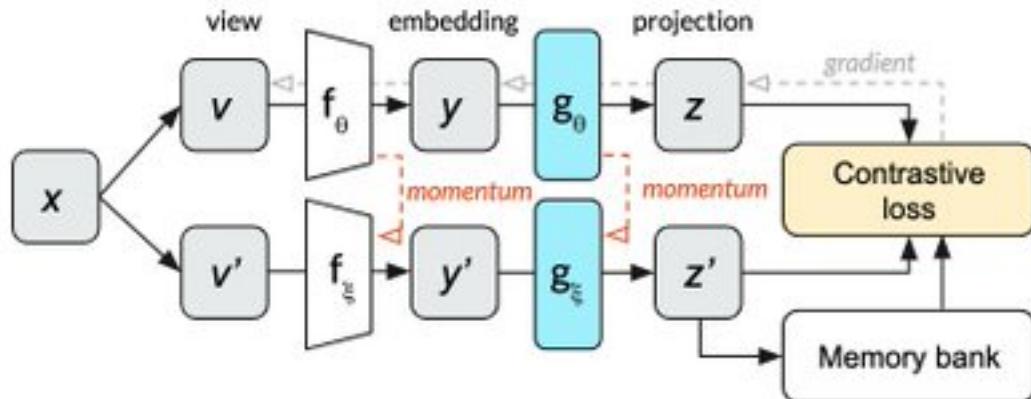
It needs a very large batch size ≥ 1024 !

The MLP that does the projection is essential.

- Learns useful transformations for the contrastive task
- But it discarded during the finetuning phase



MoCov2

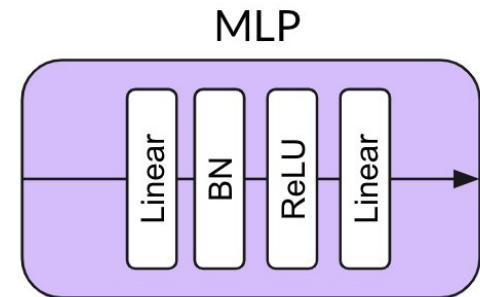
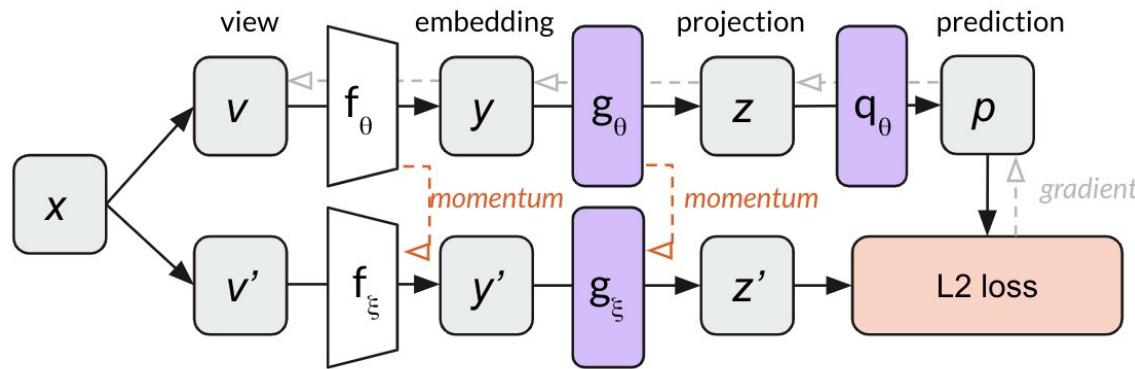


Reduces the need of large batch size with a **memory bank**:

- Stores previously computed projections z'
- Means more negative in the contrastive loss

The gradient is backpropagated only through one version of the network:

- The other network is, as in RL, a **target network**
- It is updated with momentum $\theta_t \leftarrow \alpha \theta_t + (1 - \alpha) \theta_s$
- Enforce stability in the memory bank representations



L2 distance between **only positive examples, not negative examples are used!**

Why does the representation do not collapse?

→ Meaning only producing a zero vector for any input would minimize the loss

Still an active area of research, but some intuitions:

→ Asymmetrical architecture qui another MLP q_θ

→ Momentum for the target network

What do we actually need for self-supervision?



Invariance to transformations:

- Two augmentations of the same image should produce the same representation

Disentangling of the dimensions:

- Each dimension of the representation should encode a different info

Domain Adaptation



Setting

- Source domain/dataset is fully labeled
- Target domain/dataset is unlabeled
- Both represent the same classes
- Huge discrepancy in the pixels distribution

	MNIST	SYN NUMBERS	SVHN	SYN SIGNS
SOURCE				
TARGET				
	MNIST-M	SVHN	MNIST	GTSRB

Source Domain



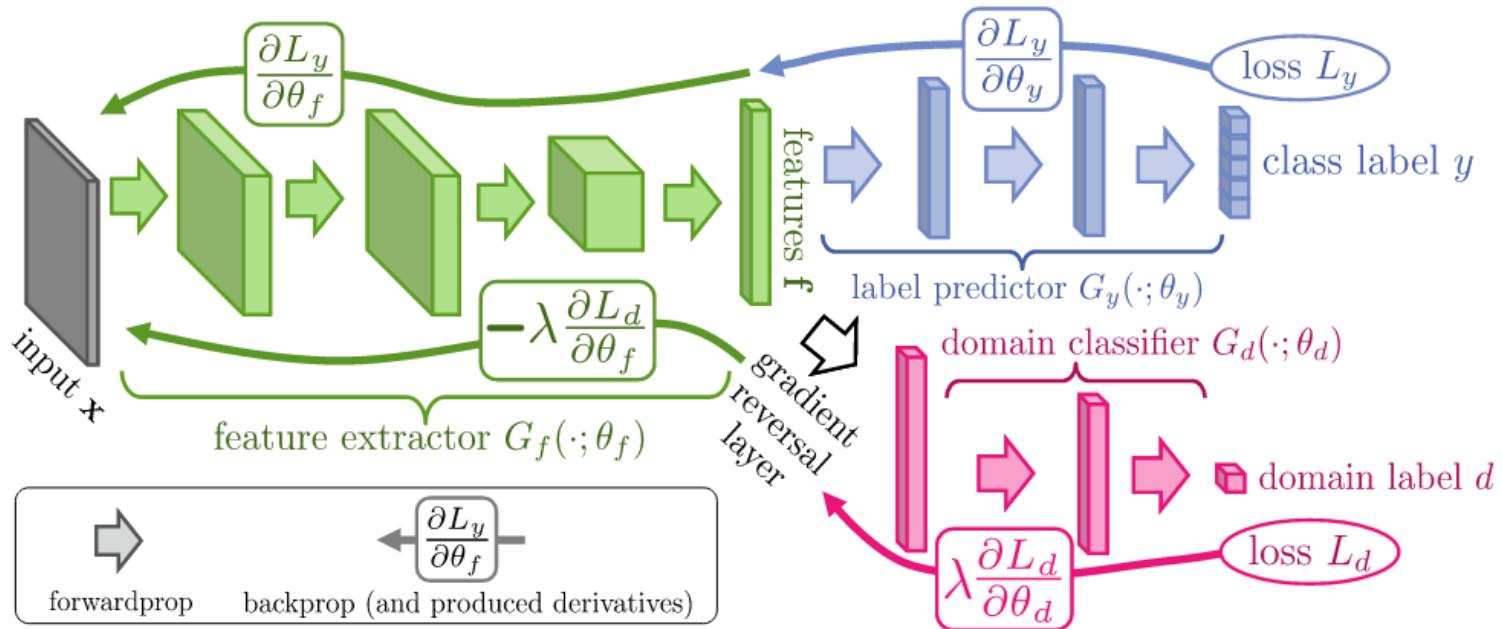
GTA5 (yes the game)

Target Domain



Cityscapes

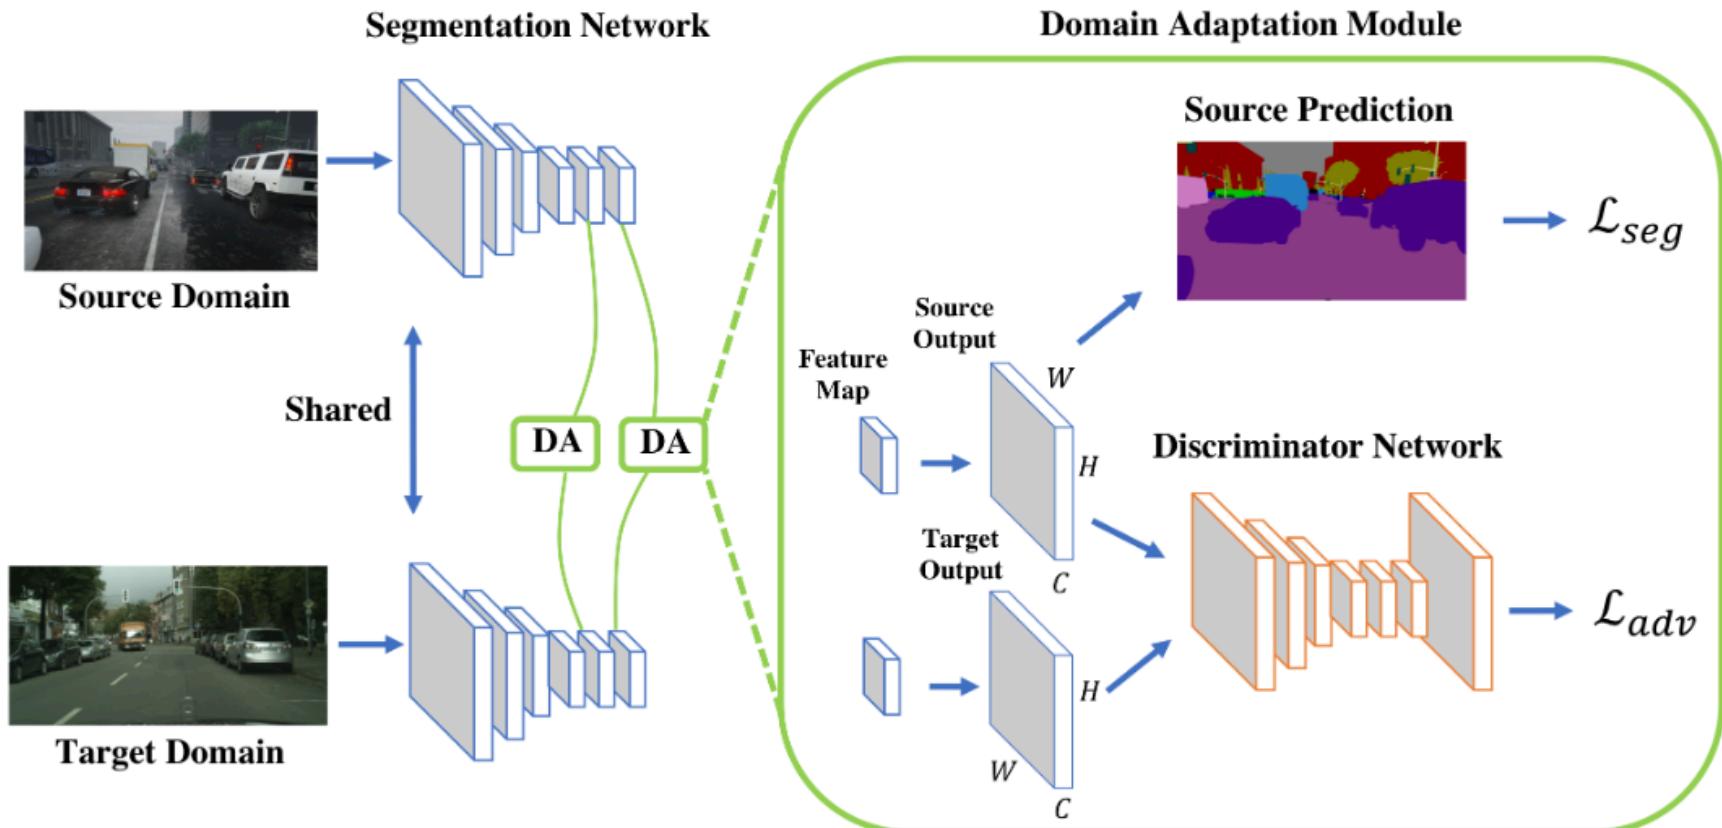
DANN: Gradient Reversal Layer



Gradient Reversal Layer (GRL) forces the **ConvNet** to maximize the loss of the **Domain Classifier**.

→ Force to learn domain agnostic features

AdaptSegNet



- Train the Discriminator on the probabilities of the source and target without the gradient flowing backward
- Train the Segmentation Network on source for classification and also force the discriminator to predict source given target images

$$\max_{\mathbf{D}} \min_{\mathbf{G}} \mathcal{L}(I_s, I_t). \quad \mathcal{L}(I_s, I_t) = \mathcal{L}_{seg}(I_s) + \lambda_{adv} \mathcal{L}_{adv}(I_t)$$



Two key ideas

Two key ideas for domain adaption in segmentation:

1. **Adversarial loss** forcing a similar representation for both source and target domains
2. **Pseudo-labeling** to generates labels for the unlabeled target domain

Other Problems

Not covered in this lecture



Zero-shot Learning:

- Not a single image of the class to predict, but access to metadata
- *Ex: understand a Wikipedia description to classify a never-seen before animal*

Semi-Supervised:

- A few amount (~10%) of the data is labeled, while the remaining is unlabeled but present

Weak supervision:

- Labels are imperfect
- *Ex: training a model to predict the hashtag on Instagram photos*

Small break,
then coding session!