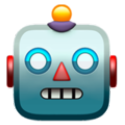

NEW ARCHITECTURES AND TRICKS TO TRAIN THEM

Deep Learning for Computer Vision

Arthur Douillard

https://arthurdouillard.com/deepcourse

Neural Architecture Search

Amoeba

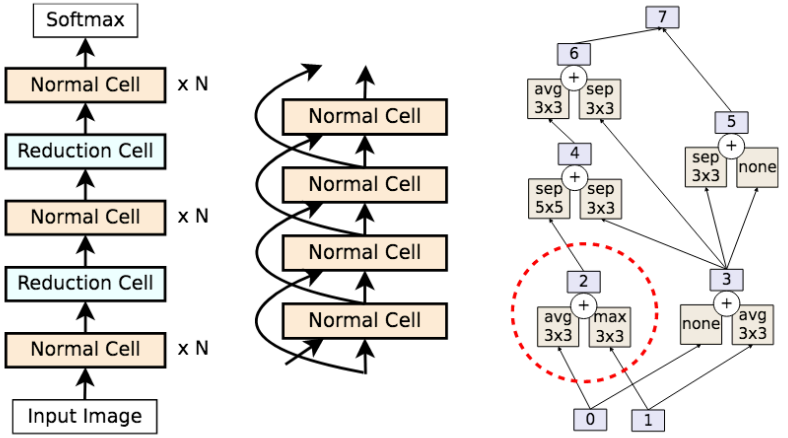

Figure 1: NASNet Search Space [54]. LEFT: the full outer structure (omitting skip inputs for clarity). MIDDLE: detailed view with the skip inputs. RIGHT: cell example. Dotted line demarcates a pairwise combination.

- Given a fixed architecture (left & middle), learn to find the optimal cell (right)
- Learning is done here with an **evolutionary algorithm** that needs to retrain & check model accuracy FOR EACH new mutation!

- NAS are usually very computation intensive, and thus it's mostly big private lab that works on it. With Quoc Le's team at Google Brain the main one.

[Real et al. AAAI 2019]

RandWire

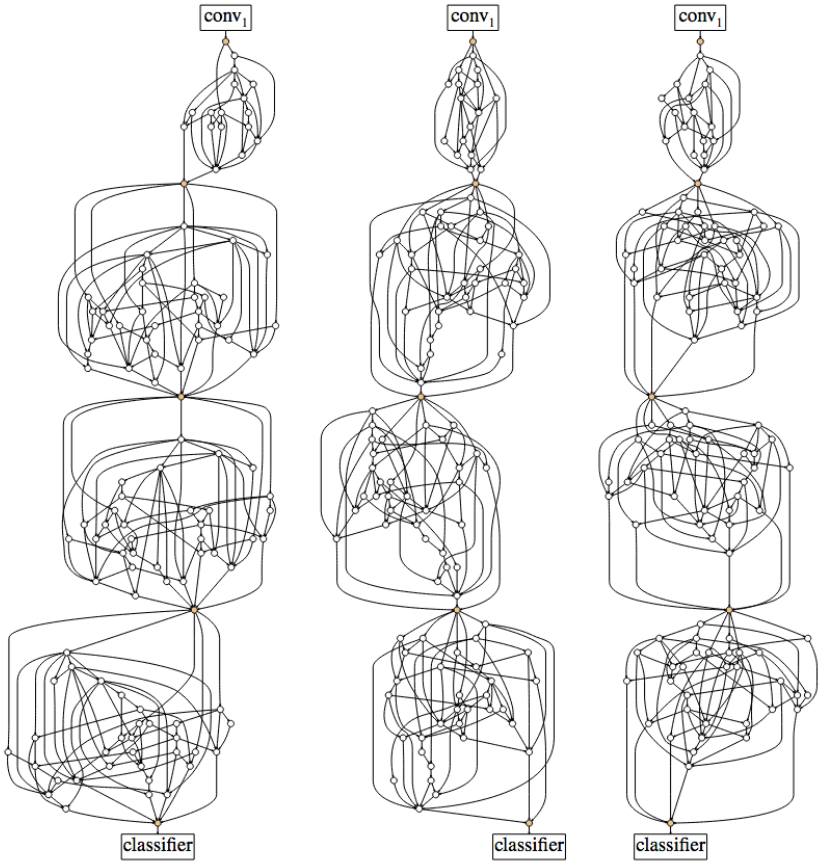

- Or try randomly wired neural network based on minimal set of rules
- But note that some approaches use reinforcement learning

[Xie et al. arXiv 2019]

EfficientNet

$$\max_{d,w,r} \quad Accuracy(\mathcal{N}(d,w,r))$$

$$s.t. \quad \mathcal{N}(d,w,r) = \bigodot_{i=1\ldots s} \hat{\mathcal{F}}_i^{d\cdot\hat{L}_i}(X_{\langle r\cdot\hat{H}_i, r\cdot\hat{W}_i, w\cdot\hat{C}_i\rangle})$$

$$\text{Memory}(\mathcal{N}) \leq \text{target\_memory}$$

$$\text{FLOPS}(\mathcal{N}) \leq \text{target\_flops}$$

$$\text{depth: } d = \alpha^\phi$$

$$\text{width: } w = \beta^\phi$$

$$\text{resolution: } r = \gamma^\phi$$

$$\text{s.t. } \alpha \cdot \beta^2 \cdot \gamma^2 \approx 2$$

$$\alpha \geq 1, \beta \geq 1, \gamma \geq 1$$

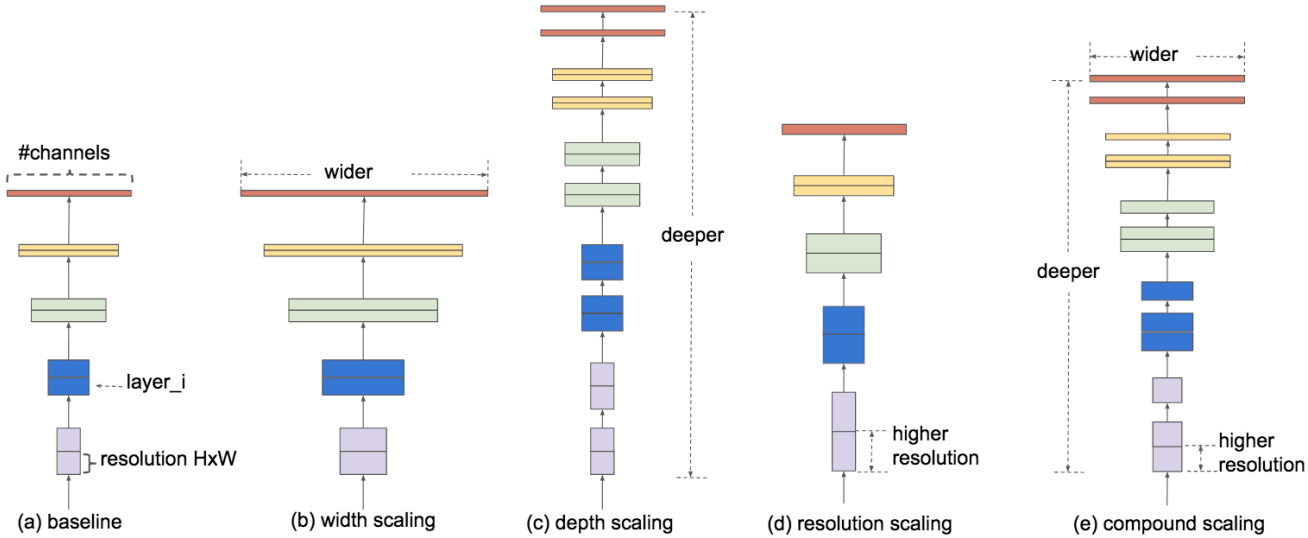

- EfficientNet, one of the best ConvNet as of 2021, was made with NAS
- Based on a **compound scaling** rule they drastically reduce the space to grid search

[Tan and Le, ICML 2019]

Transformers

Attention is all you need

- At first designed for machine translation
- Does not rely on Conv1d or RNN

- Main block are FC layers and the famous **Multi-Head Attention**

- Made of a encoder (left) and decoder (right)

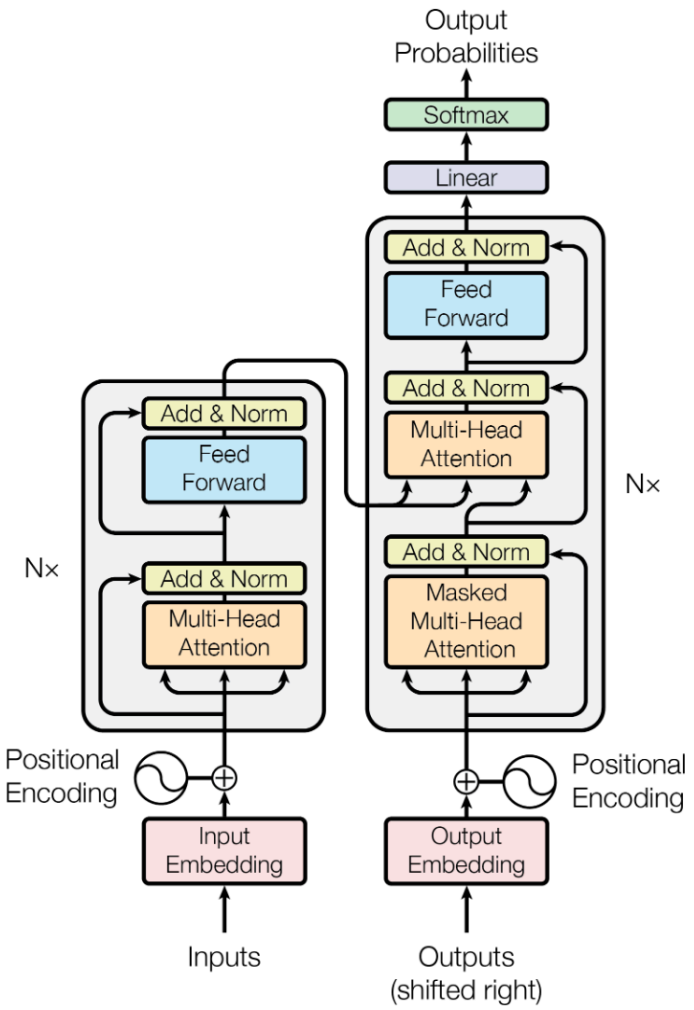

[Tan and Le, ICML 2019]

BERT

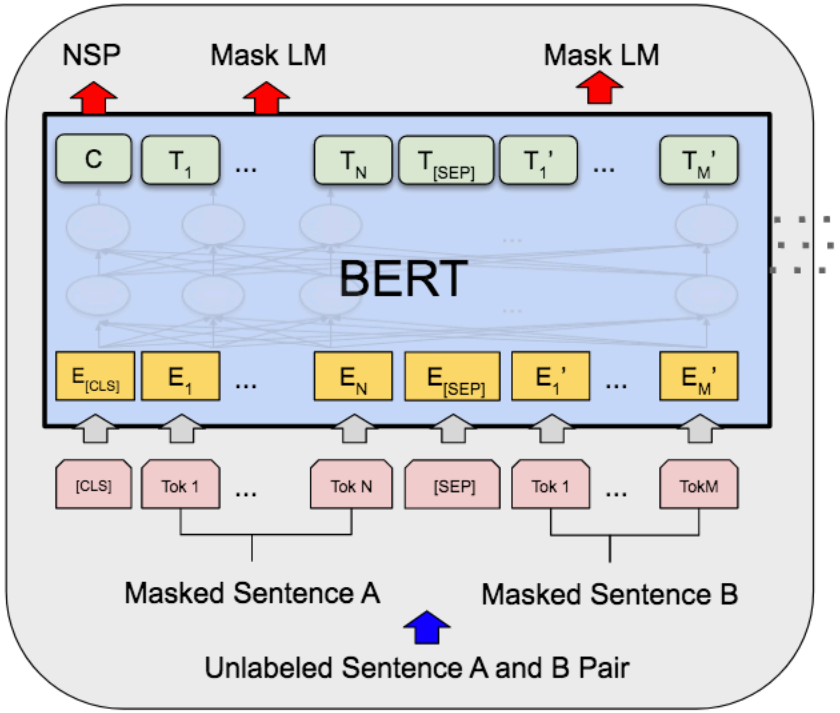

- Modification of the Transformer
- No more encoder/decoder, only many blocks
- Introduction of a special token [CLS] that is learned
- Was a big revolution in the NLP field

[Devlin et al. NAACL 2019]

ViT

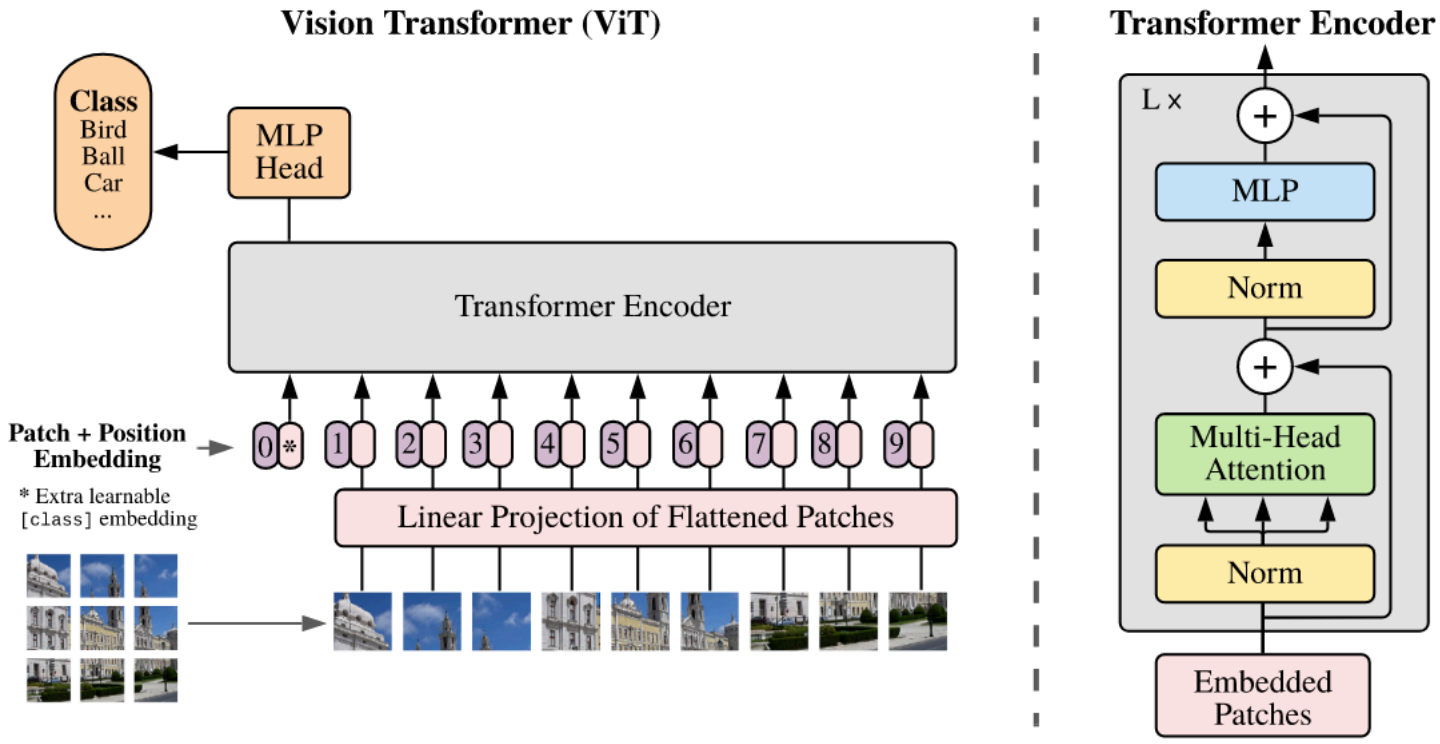

- First *full* application of transformers with a BERT-like architecture to vision

[Dosovitskiy, Beyer, Kolesnikov et al. ICLR 2021]

Patch, Position, and class token

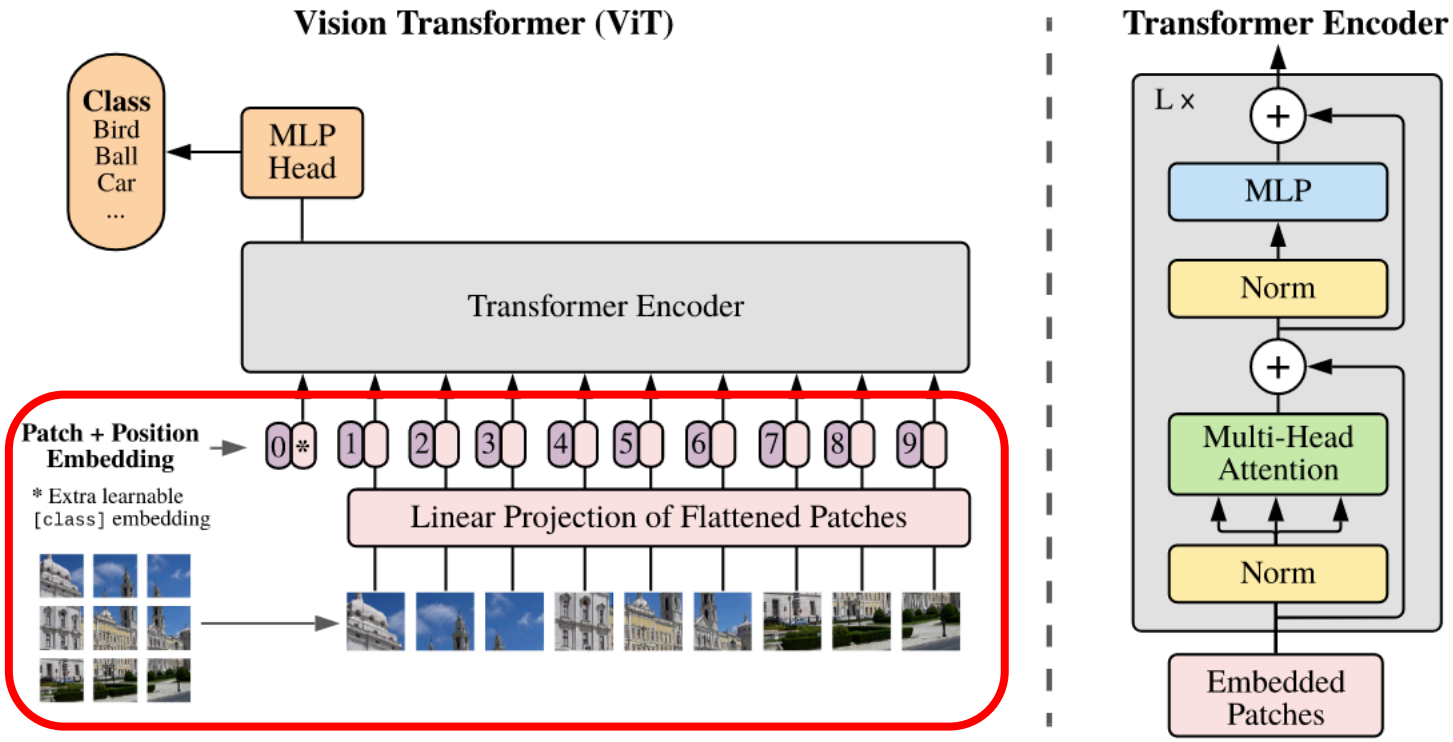

- Use a 1x1 convolution with a kernel size equal to the patch size to generate the **tokens**
  - Total size is thus *(batch size, number of tokens, embedding dimension)*
- Add an **extra token [class]** that is a learned vector of size *(embedding dimension)*
- Add to all tokens a learned **positional embeddings**

[Dosovitskiy, Beyer, Kolesnikov et al. ICLR 2021]

Self-Attention

1. Apply three different linear transformations, to create the **Q**uery, the **K**ey, and the **V**alue

$$Q = XW_q$$

$$K = XW_k$$

$$V = XW_v$$

2. Compute the attention matrix that measures the inter-tokens similarity

$$A = \text{softmax}(\frac{QK^T}{\sqrt{d}})$$

3. Ponderate the **V**alue matrix with the attention matrix

$$Z = AV$$

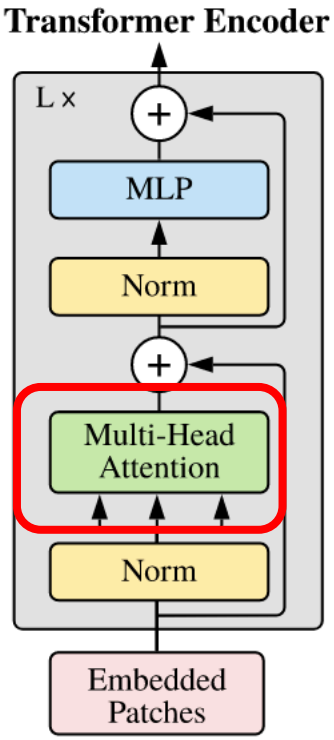

[Dosovitskiy, Beyer, Kolesnikov et al. ICLR 2021]

Multi-heads Self-Attention

$$Q_1 = XW_{q_1} \text{ and } Q_2 = XW_{q_2}$$

$$K_1 = XW_{k_1} \text{ and } K_2 = XW_{k_2}$$

$$V_1 = XW_{v_1} \text{ and } V_2 = XW_{v_2}$$

$$Z_1 = \text{softmax}(\frac{Q_1 K_1^T}{\sqrt{d}}) \text{ and } Z_2 = \text{softmax}(\frac{Q_2 K_2^T}{\sqrt{d}})$$

$$Z' = [Z_1 Z_2] \in \mathbb{R}^{T \times 2d}$$

$$Z = Z' W_o$$

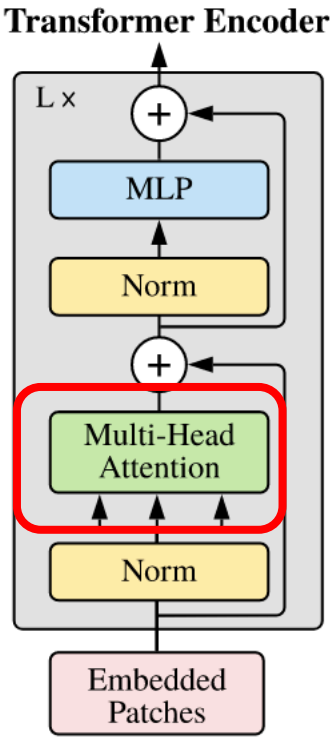

In practice, the self-attention is done multiple times in parallel, with different **heads**.

[Dosovitskiy, Beyer, Kolesnikov et al. ICLR 2021]

DeiT

Ablation on ↓	Pre-training	Fine-tuning	Rand-Augment	AutoAug	Mixup	CutMix	Erasing	Stoch. Depth	Repeated Aug.	Dropout	Exp. Moving Avg.	top-1 accuracy: pre-trained 224^2	top-1 accuracy: fine-tuned 384^2
none: DeiT-B	adamw	adamw	✓	✗	✓	✓	✓	✓	✓	✗	✗	81.8 ±0.2	83.1 ±0.1
optimizer	SGD	adamw	✓	✗	✓	✓	✓	✓	✓	✗	✗	74.5	77.3
	adamw	SGD	✓	✗	✓	✓	✓	✓	✓	✗	✗	81.8	83.1
data augmentation	adamw	adamw	✗	✗	✓	✓	✓	✓	✓	✗	✗	79.6	80.4
	adamw	adamw	✗	✓	✓	✓	✓	✓	✓	✗	✗	81.2	81.9
	adamw	adamw	✓	✗	✗	✓	✓	✓	✓	✗	✗	78.7	79.8
	adamw	adamw	✓	✗	✓	✗	✓	✓	✓	✗	✗	80.0	80.6
	adamw	adamw	✓	✗	✗	✗	✓	✓	✓	✗	✗	75.8	76.7
regularization	adamw	adamw	✓	✗	✓	✓	✗	✓	✓	✗	✗	4.3*	0.1
	adamw	adamw	✓	✗	✓	✓	✓	✗	✓	✗	✗	3.4*	0.1
	adamw	adamw	✓	✗	✓	✓	✓	✓	✗	✗	✗	76.5	77.4
	adamw	adamw	✓	✗	✓	✓	✓	✓	✓	✓	✗	81.3	83.1
	adamw	adamw	✓	✗	✓	✓	✓	✓	✓	✗	✓	81.9	83.1

- Transformers are hard to train because they lack **inductive bias**, and thus needs way more data than a ConvNet
- DeiT partially close this gap by using tons of **data augmentations** and **regularizations**

[Touvron et al. ICML 2021]

CaiT

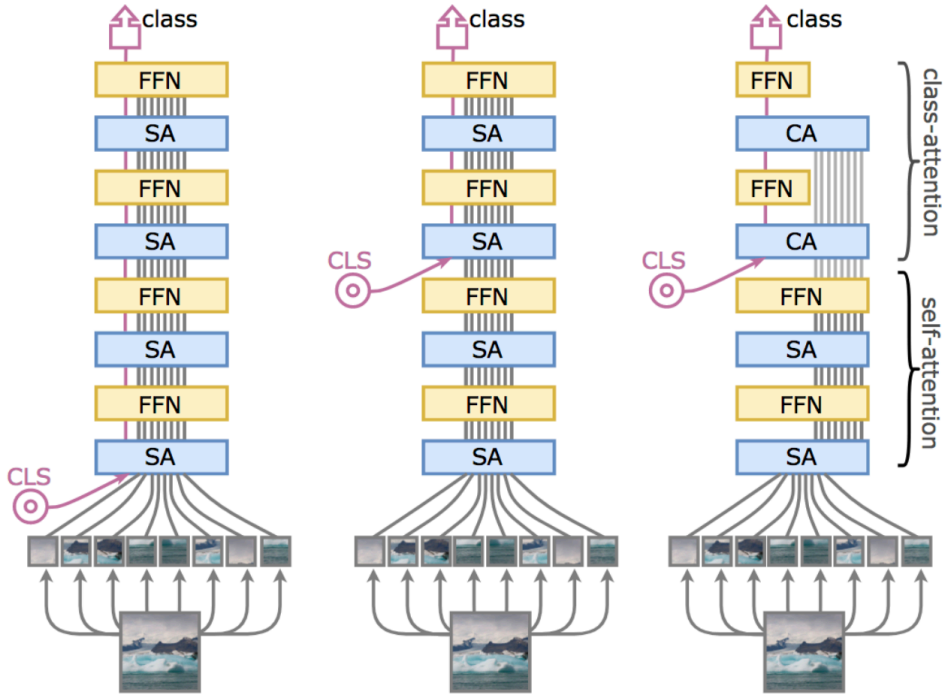

- Inserting the class tokens in later blocks may prove beneficial

[Touvron et al. ICCV 2021]

Future of Transformers

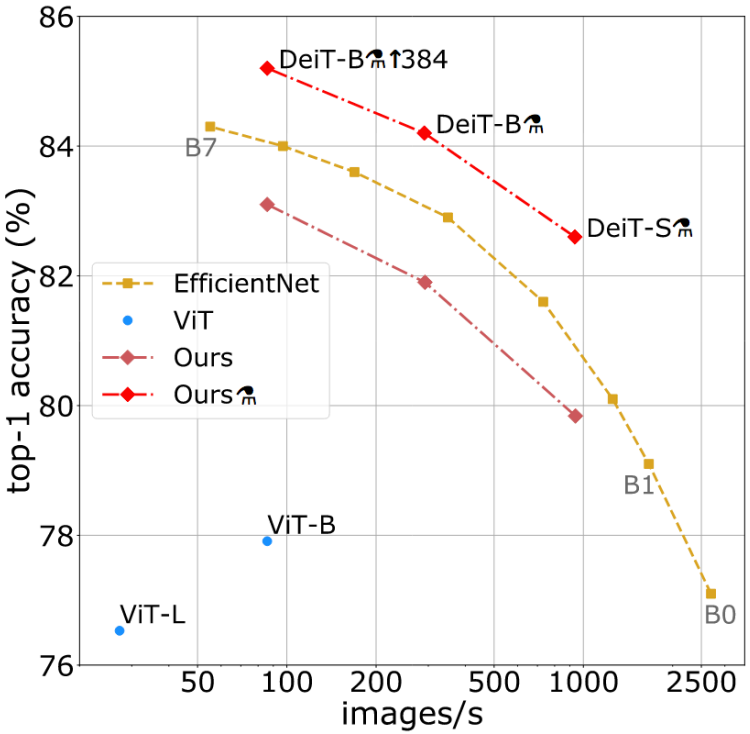

- More generally transformers for vision seems to be leading the SotA in computer vision
- And they manage to be super fast

- Although, they need more data and regularizations than your common ConvNet
- And using loss curvature regularization (like SAM) seems to help a lot

- Hard to follow literature, hundred of vision transformers in Jan—Aug 2021! Time will tell which ones are really useful

Fourier Transformer

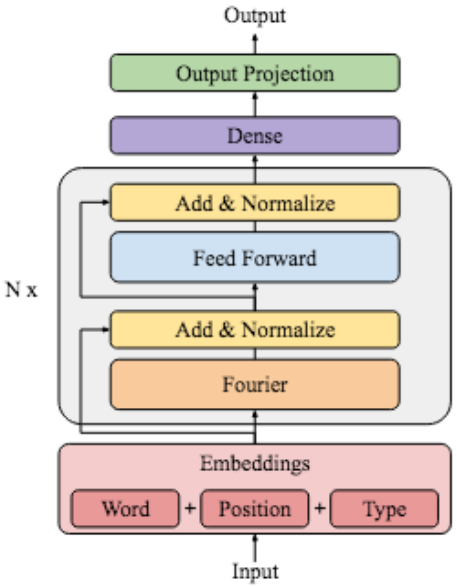

Fourier transform that can be expressed as a matrix multiplication with this constant matrix:

$$W_{nk} = \left(e^{-\frac{2\pi i}{N} nk} / \sqrt{N} \right)$$

- Not as good as Self-Attention, but still impressive results

- Means that ”*Attention is **NOT** all you need*”, but rather a way to combine inputs
  - Likewise convolutions combine pixels through the increasing receptive field

MLP Comeback!

Separable Convolutions

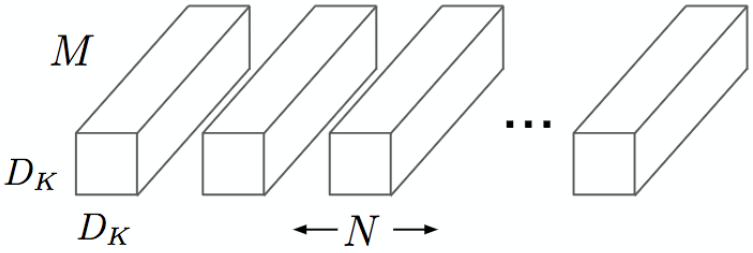

(a) Standard Convolution Filters

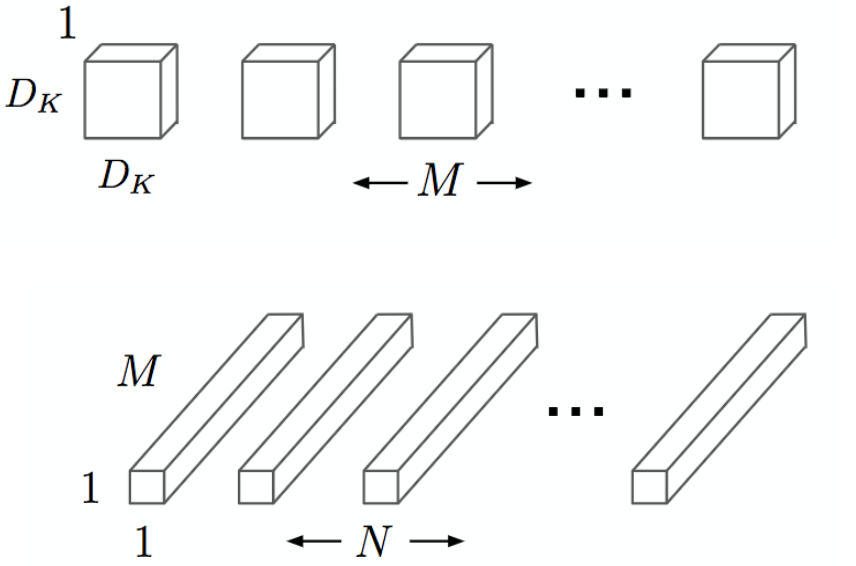

(c) 1×1 Convolutional Filters called Pointwise Convolution in the context of Depthwise Separable Convolution

- **Depthwise convolutions**: doesn't mix input channels
- **Pointwise convolutions**: doesn't mix spatial dimensions

[Chollet, arXiv 2016]

MLP-Mixer

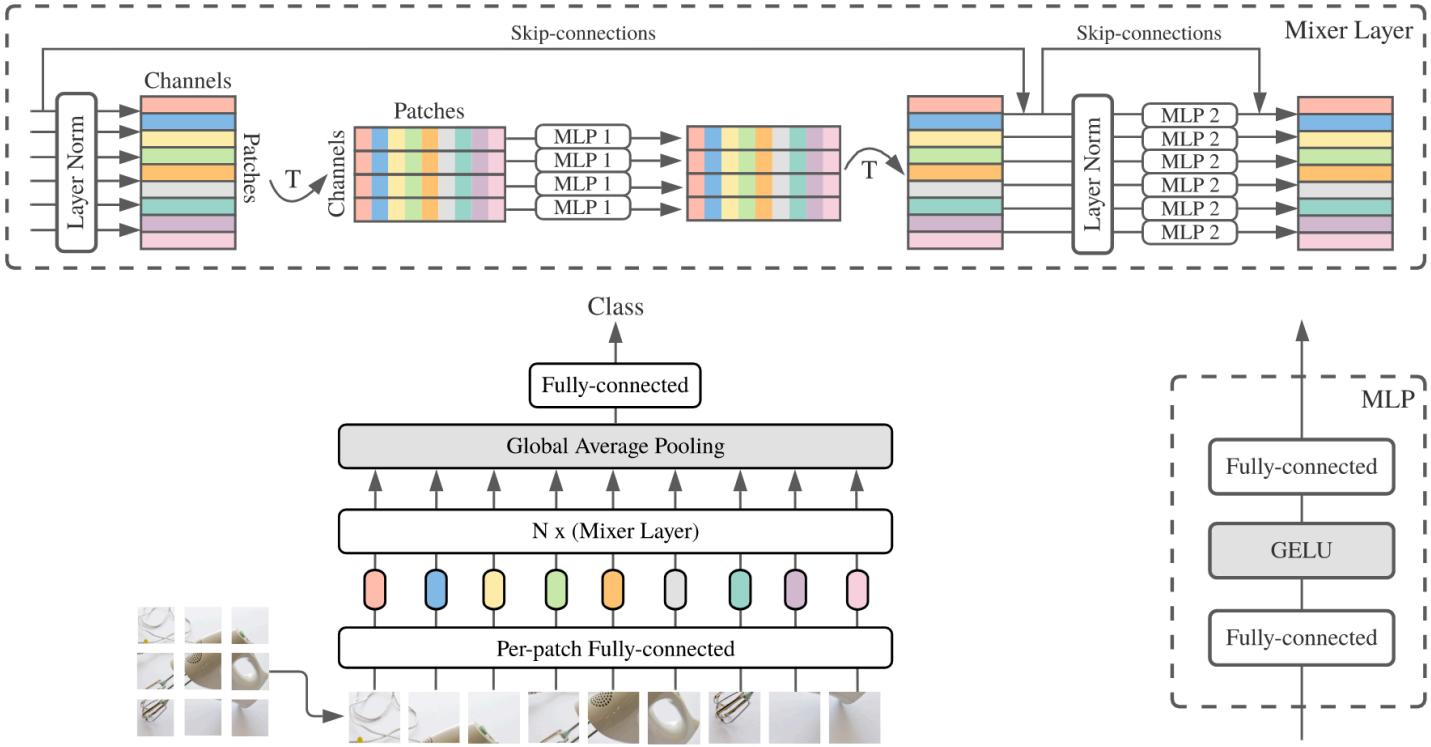

- Similarly to Separable Convolutions, apply a MLP on the channels dimension and a MLP on the patches dimensions

- Multiple other papers had the same idea at the same time (including ResMLP)

[Tolstikhin and Houlsby and Kolesnikov and Beyer et al. arXiv 2021]

MLP-Mixer

	ImNet top-1	ReaL top-1	Avg 5 top-1	VTAB-1k 19 tasks	Throughput img/sec/core	TPUv3 core-days
Pre-trained on ImageNet-21k (public)						
HaloNet [51]	85.8	—	—	—	120	0.10k
Mixer-L/16	84.15	87.86	93.91	74.95	105	0.41k
ViT-L/16 [14]	85.30	88.62	94.39	72.72	32	0.18k
BiT-R152x4 [22]	85.39	—	94.04	70.64	26	0.94k
Pre-trained on JFT-300M (proprietary)						
NFNet-F4+ [7]	89.2	—	—	—	46	1.86k
Mixer-H/14	87.94	90.18	95.71	75.33	40	1.01k
BiT-R152x4 [22]	87.54	90.54	95.33	76.29	26	9.90k
ViT-H/14 [14]	88.55	90.72	95.97	77.63	15	2.30k
Pre-trained on unlabelled or weakly labelled data (proprietary)						
MPL [34]	90.0	91.12	—	—	—	20.48k
ALIGN [21]	88.64	—	—	79.99	15	14.82k

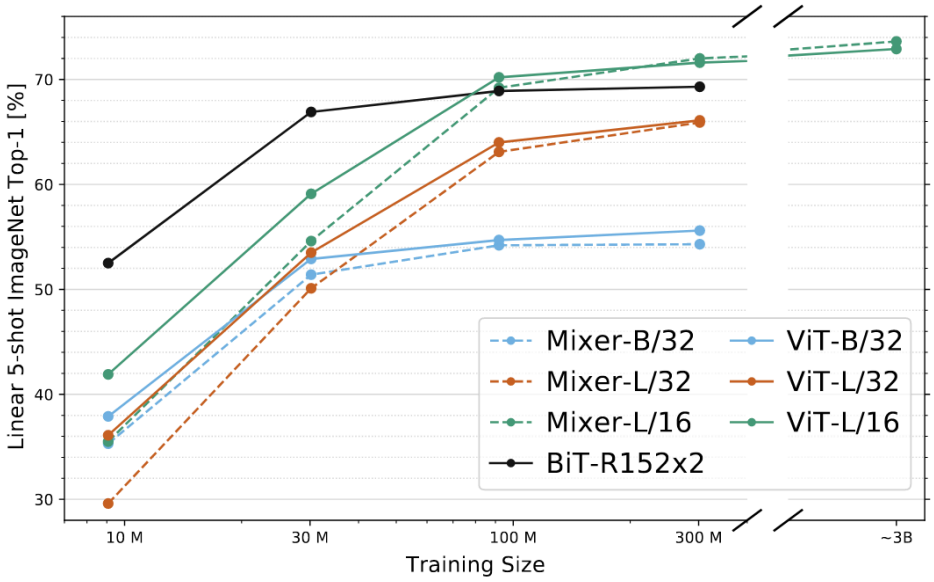

- However needs even more data than transformers based (ViT)
- Convolutions models (here BiT-R152x2) are still much more data efficient when training from scratch
  - However, some advocates that most future applications will be based on those new large (transformer, mlp, etc.) models pretrained on large-scale data
  - See [Bommasani et al. arXiv 2021]

Tricks that work for most architectures

Train / Val / Test

- It’s extremely important to tune the hyperparameters on a val set

- In real-life, you often don’t have access to the full test set
  - And this test set may change constantly

- It’s also important to ensure that your train and val sets have the same distribution than the test set
  - Beware of the **sampling bias**, e.g. I’m only labeling images of cars oriented towards the front, but in the test / real-life I may see cars in other orientations
  - See [Torralba and Efros, CVPR 2011]

Active Learning

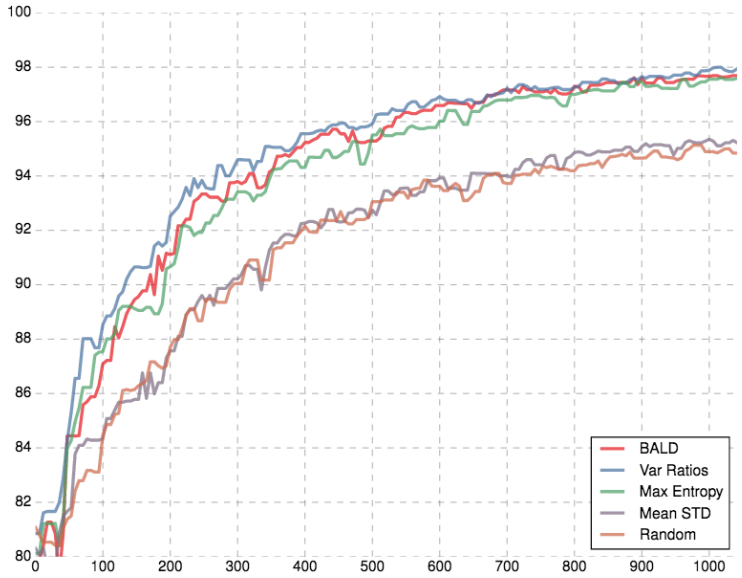

- More data is (almost) always better, see the Sutton's Bitter Lesson

- But hand-labeling data is very costly, both in $ and time

- Active learning aims to determine which data to labelize in priority to be added to the training set
  - A lot of the literature is based on Bayesian stats, with Yarin Gal's team
  - But often done on small-scale datasets (MNIST, CIFAR)
  - And a random sampling is often quite competitive despite its simplicity

[Gal et al. ICML 2017]

Thresholding for open-set prediction recall/precision/f1

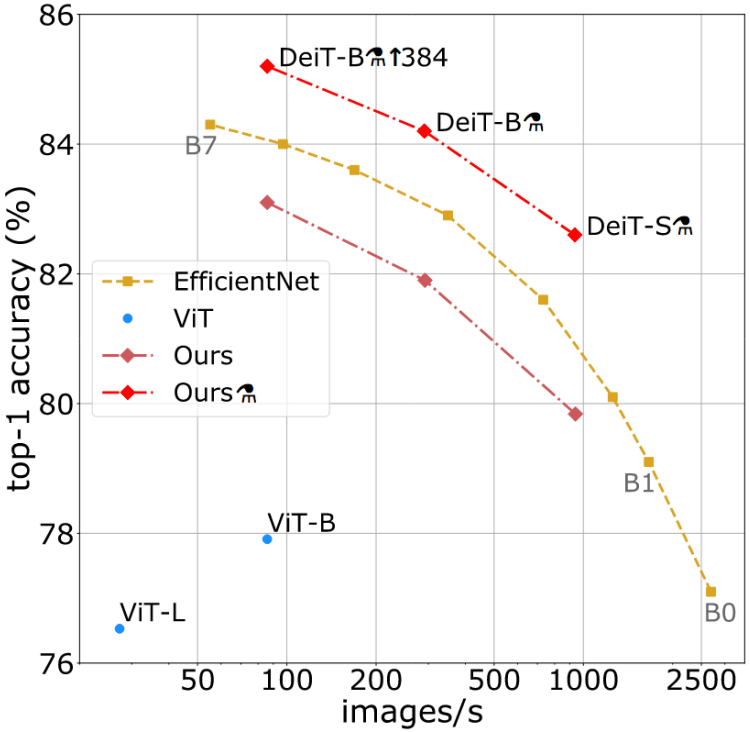

- More generally transformers for vision seems to be leading the SotA in computer vision
- And they manage to be super fast

- Although, they need more data and regularizations than your common ConvNet

[Touvron et al. ICML 2021]

Early Stopping

- Training for too long, and you start to overfit. So when to stop?

- Monitor a metric (accuracy, loss, etc.) on the **VALIDATION SET** (!) and if this metric gets worse for X epochs, stop training

- ”*A beautiful free lunch*” according to Turing award’s Geoffrey Hinton

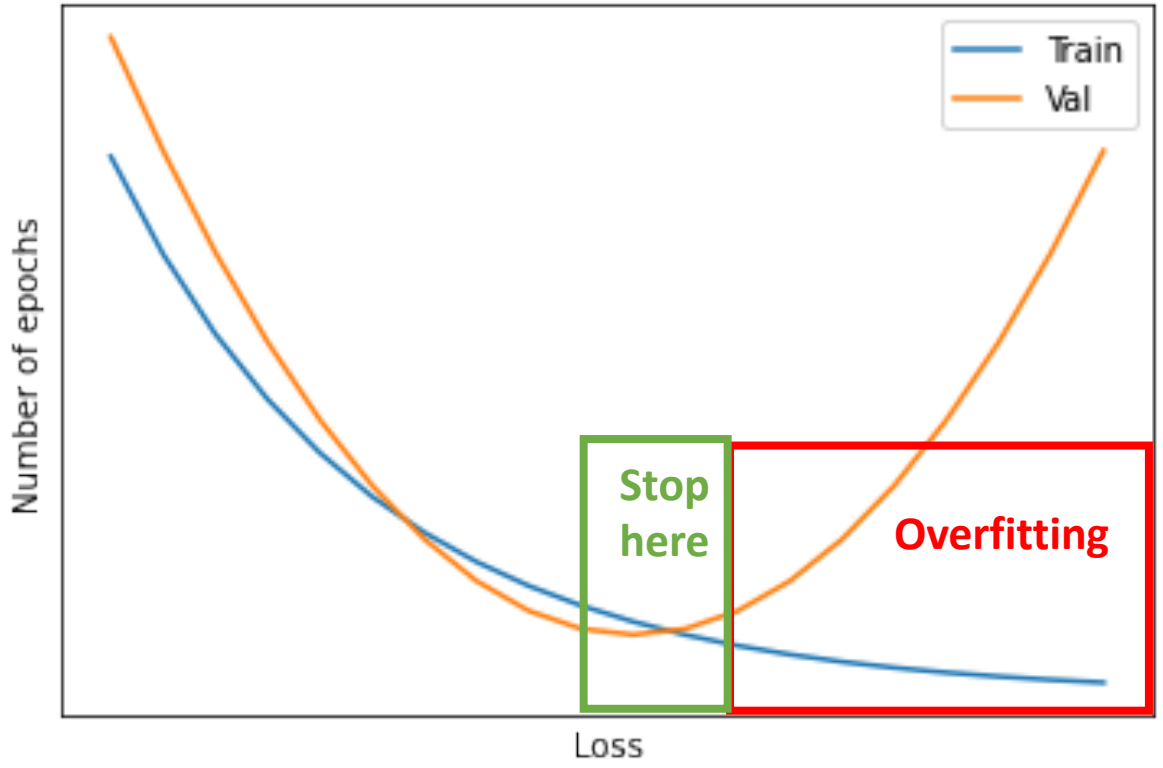

Decrease Learning Rate

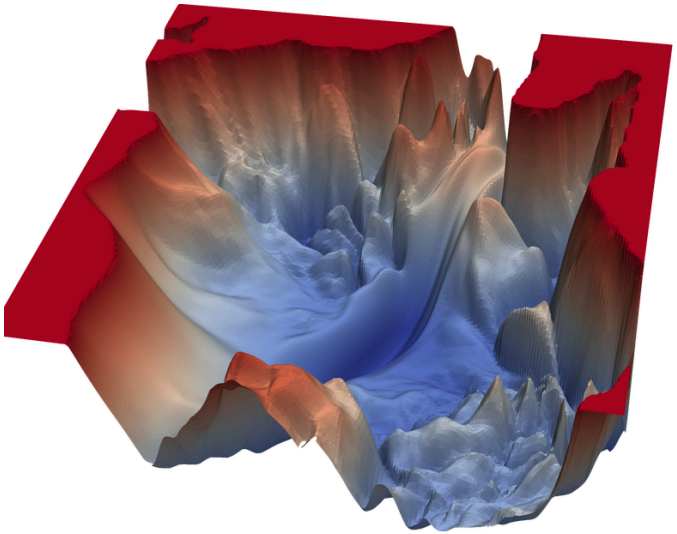

- A high learning rate during the beginning of the training may help
  - Acts as a regularization by **skipping the local minima that are too sharp** and thus usually generalize less

- Then **decrease learning rate gradually to go deeper in a local minim**a towards the training end
  - Either decrease learning rate (usually divided by 10) at particular epochs
  - Or decrease if validation metric (loss, acc, etc.) doesn't improve

- Some scheduling as **Cosine** decreases and increases (a little less) repetitively

Smooth Labeling

$$y_k^{LS} = y_k(1-\alpha) + \alpha/K$$

[0, 0, 1, 0] with a $\alpha = 0.1$ produces [0.025, 0.025, 1, 0.025]

- Avoid overconfidence in the model, when confidence is always close to 0.999…
  - And thus reduce overfitting

- Avoid miscalibration where model confidence is not correlated to model accuracy

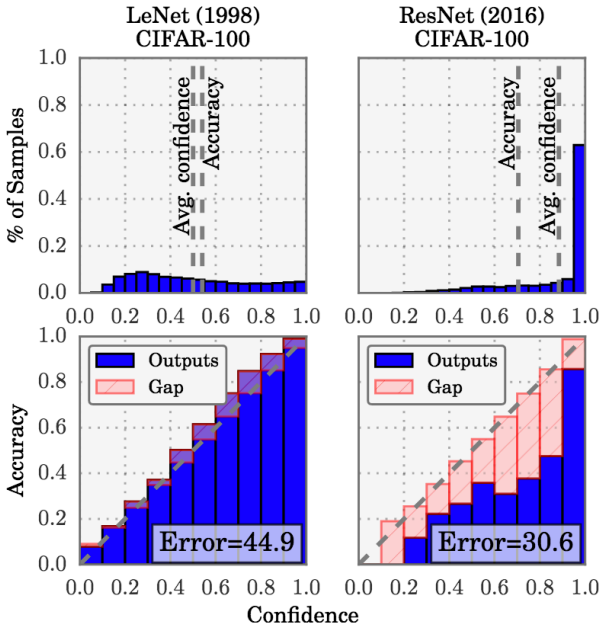

Smooth labeling: [Muller et al. NeurIPS 2019], Miscalibration: [Guo et al. ICML 2017]

MixUp & Co

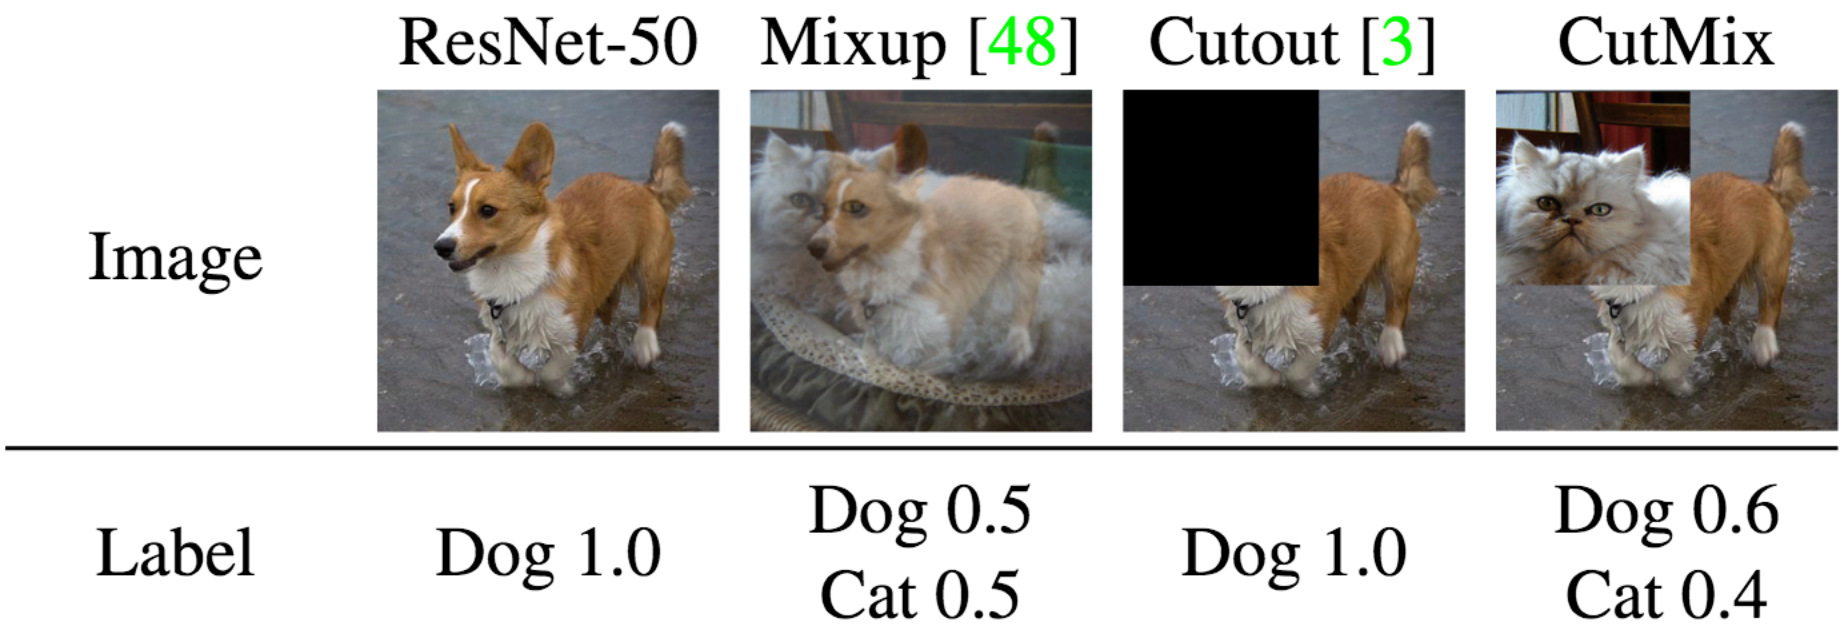

- Mix two images and their labels together
  - In practice MixUp mix them with factors like 0.9/0.1 (not 0.5/0.5 as on the image)
- Acts as regularization to reduce overfitting
- A LOT of alternative to MixUp exists (CutOut, CutMix, FixMatch, MixMo, PuzzleMix, etc.)

Mixup: [Zhang et al. ICLR 2017]

Knowledge Distillation

- Train one super-mega-large model (called teacher)
- **Distill** the knowledge of the teacher onto a smaller model (called student)

- In practice, train students as usual but add a another loss:
  - **KL-divergence** between the probabilities of the teacher and the student
  - The probabilities acts as **dark knowledge** with extra information
    - (aka if the teacher say this is a dog with 0.7 confidence, we know it's a dog, but it's probably not the most archetypal dog ever)
  - Often add a **temperature** $\boldsymbol{T}$ on the logits before softmax
    - If $T > 1$, reduces the sharpness of the probabilities leading too more useful info

$$q_i = \frac{exp(z_i/T)}{\sum_j exp(z_j/T)}$$

[Hinton et al. NeurIPS Workshop 2014]

Thresholding & Open-Set

- In the real-life test set, you don’t always want to predict on all images

- Example:
  - *at Heuritech we predict fashion trends from images scrapped from Instagram. Given an image of dog, my model should predict any trend!*

- You choose to discard model prediction if its confidence is lower than a threshold
  - One threshold per class is better
  - Compute threshold on the validation set (which needs to have negative images!)

Accuracy vs Precision+Recall=F1

- In real-life classes are never equally balanced
- If there are 90% of *dogs*, and 10% of *cats*, and your model has less than 90% of accuracy it's bad…
  - Answering *dog* every times gives 90% accuracy
- Best to use other metrics like **Precision** and **Recall** or their combination the **F1-Score**
- **Recall** is particularly useful in **open-set** where your model shouldn't predict on all images

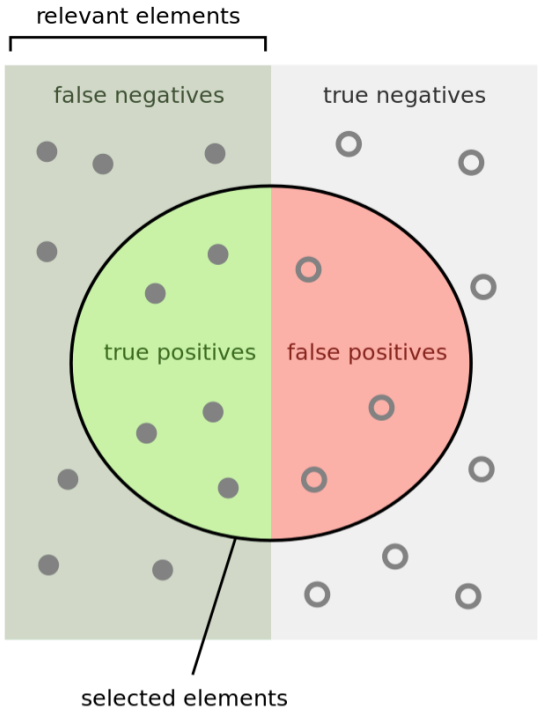

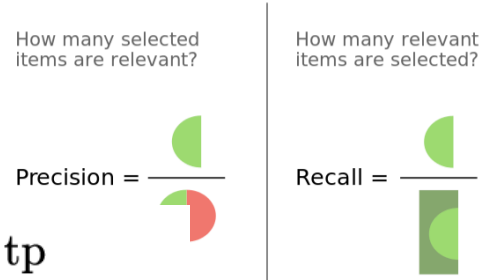

$$F_1 = \frac{2}{\mathrm{recall}^{-1} + \mathrm{precision}^{-1}} = 2 \cdot \frac{\mathrm{precision} \cdot \mathrm{recall}}{\mathrm{precision} + \mathrm{recall}} = \frac{\mathrm{tp}}{\mathrm{tp} + \frac{1}{2}(\mathrm{fp} + \mathrm{fn})}$$

Wikipedia

SAM: Sharpness-Aware Minimization

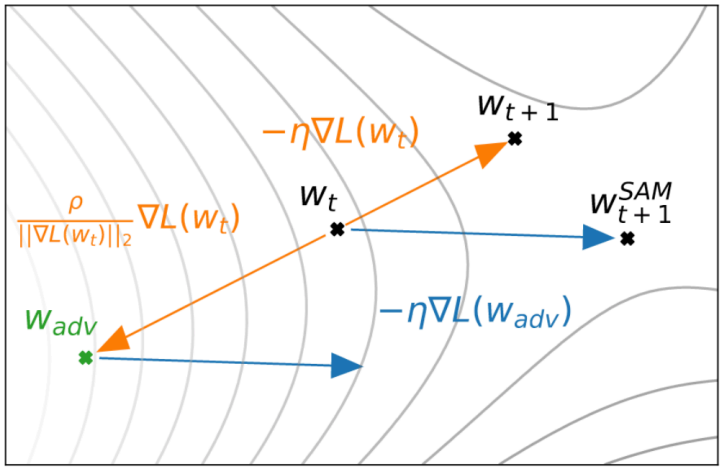

Input: Training set $\mathcal{S} \triangleq \cup_{i=1}^{n}\{(\boldsymbol{x}_i, \boldsymbol{y}_i)\}$, Loss function $l : \mathcal{W} \times \mathcal{X} \times \mathcal{Y} \rightarrow \mathbb{R}_+$, Batch size b, Step size $\eta > 0$, Neighborhood size $\rho > 0$.
Output: Model trained with SAM
Initialize weights $\boldsymbol{w}_0$, $t = 0$;
while *not converged* **do**
  Sample batch $\mathcal{B} = \{(\boldsymbol{x}_1, \boldsymbol{y}_1), ...(\boldsymbol{x}_b, \boldsymbol{y}_b)\}$;
  Compute gradient $\nabla_{\boldsymbol{w}} L_{\mathcal{B}}(\boldsymbol{w})$ of the batch's training loss;
  Compute $\hat{\boldsymbol{\epsilon}}(\boldsymbol{w})$ per equation 2;
  Compute gradient approximation for the SAM objective (equation 3): $\boldsymbol{g} = \nabla_w L_{\mathcal{B}}(\boldsymbol{w})|_{\boldsymbol{w}+\hat{\boldsymbol{\epsilon}}(\boldsymbol{w})}$;
  Update weights: $\boldsymbol{w}_{t+1} = \boldsymbol{w}_t - \eta \boldsymbol{g}$;
  $t = t + 1$;
end
return $\boldsymbol{w}_t$

- Do not optimize network on a particular point of the parameters space but rather a **region**
- All neighbors parameters must also be good, leading to wider optimum and thus better generalization
- Needs twice more forward/backward…

[Foret et al. ICLR 2021]

Small break,
then coding session!